



UNIVERSITÀ DI PISA

Department of Computer Science
Bachelor's Degree in Computer Science

BACHELOR'S THESIS

**Can we recover differential transcript
usage (DTU) from scRNA-seq?**

Supervisors:
Prof.ssa Silvia Galfré
Prof. Corrado Priami

Candidate:
Lorenzo Delitala

Academic Year 2025/2026

Abstract

Single-cell RNA sequencing has transformed the study of cellular heterogeneity, yet widely adopted droplet-based protocols capture only the 3' end of each transcript [2]. This constraint makes Differential Transcript Usage (DTU) difficult to detect, as short terminal reads lack the internal splice junctions required to distinguish isoforms. Furthermore, proportion-based DTU methods—including those extended to single-cell data, which rely on pseudo-bulk aggregation—sacrifice cell-level resolution to recover the read depth they require [3, 4].

This thesis investigates whether COTAN—a statistical framework that models zero counts directly through presence/absence contingency tables [1]—can be extended from gene-level co-expression to transcript-level DTU detection. By bypassing both artificial imputation and pseudo-bulk aggregation, the adapted method identifies mutually exclusive isoform pairs via independence testing and bidirectional contrast thresholds across cell clusters.

A reproducible Nextflow pipeline was built to construct transcript-level matrices from raw 3'-biased data using alevin-fry with identity transcript-to-gene mapping [5]. This pipeline was applied to a human cell line mixture (29,000 cells) and mouse brain tissue (4,000 cells). Transcript-level clustering successfully recovered fine-grained substructures consistent with the Global Differentiation Index, and top DTU candidates—including *Meg3*, *Malat1*, and *TPM1*—were consistent with documented isoform biology [6].

These results demonstrate that sparsity-robust, contingency-table statistics can detect isoform mutual exclusivity in droplet-based data without the prohibitive cost of full-length sequencing. While sensitivity remains bounded by positional bias—rendering mid-transcript splicing events invisible—and orthogonal validation is still required, this work establishes a principled path for isoform-aware analysis on large-scale atlas datasets.

Contents

Abstract	iii
Glossary	vii
1 Introduction	1
1.1 Biological Context for Computer Scientists	1
1.2 A Brief History of Single-Cell Sequencing	2
1.3 The Computational Challenge	3
1.4 Current Approaches and Limitations	4
1.5 Thesis Objective	4
2 Theory and Frameworks	7
Notation	7
2.1 Preprocessing Pipeline	8
2.2 COTAN Framework	16
2.3 DTU Fundamentals	20
3 Materials and Methods	23
3.1 Data Acquisition and Dataset Selection	23
3.2 Alignment, Quantification, and Feature Generation	24
3.3 Data Processing and Quality Control	27

3.4	Population Partitioning and Validation	28
3.5	The DTU Detection Algorithm	29
3.6	Calibrating the Switch Threshold (τ)	30
4	Results and Discussion	31
4.1	Dataset Quality Assessment via GDI	31
4.2	Clustering Comparison: Gene vs Transcript Level	32
4.3	Core Finding: DTU Detection Results	34
4.4	Discussion	38
5	Conclusions	41
5.1	Summary of Findings	41
5.2	Broader Implications	41
5.3	What This Work Leaves Behind	42
5.4	Limitations and Future Work	43
	Bibliography	45
A	Supplementary Data	49
A.1	Complete DTU Results (Arigoni and Ding Datasets)	49

Glossary

This glossary defines the core concepts whose discussion spans several chapters. Terms needed only once are defined inline as footnotes where they first appear; the entries below are those the reader may want to revisit. Each term is defined here and hyperlinked from its first mention in the text.

3' end The terminal end of an RNA molecule at which transcription finishes; in mRNA, it is where the poly-A tail is added. In 3'-biased protocols, the capture chemistry is anchored at this end, so reads cover only the terminal fragment of each molecule rather than the whole transcript. 3, 38

3'-biased sequencing A property of some scRNA-seq protocols (such as 10x Genomics Chromium) in which sequencing reads are derived only from the 3' terminus of each transcript. Because reading quality decreases by distance and because capture is anchored at the transcript's 3' end, coverage does not extend into the transcript body, so internal splice junctions are never observed [7]. 4, 13, 14

alternative splicing The process by which a gene's long primary RNA copy is shortened and rearranged during maturation. The copy contains coding blocks (exons) and non-coding blocks (introns); splicing removes the introns and joins the flanking exons at the resulting splice junctions. When exon selection varies between transcripts of the same gene, the result is multiple distinct isoforms. 1, 20

bulk RNA-seq Sequencing of RNA pooled from many cells at once, which averages away cell-to-cell differences. 2, 4, 17

cell barcode A short DNA tag carried by every read from a given cell, identifying its cell of origin. It allows the reads of one pooled sequencing run to be regrouped into per-cell expression profiles. 2, 10, 15

co-expression The tendency of two features (genes or transcripts) to be detected together across cells. 27, 29

contingency table A 2×2 table counting, across all cells, the four presence/absence combinations of two features; it tests whether they co-occur more or less than chance. 17, 29

count matrix A table whose rows index features (genes or transcripts) and whose columns index cells, with each entry recording how many molecules of that feature were detected in that cell. 3, 4, 8, 23

differential gene expression (DGE) The analysis of differences in a gene’s total expression—the summed expression of all its isoforms—between groups of samples. It measures overall gene output and is blind to changes in the relative usage of individual isoforms, which are instead captured by differential transcript usage. 3,

20

Differential Transcript Usage (DTU) Differential transcript usage occurs when the relative expression of multiple transcripts arising from the same gene changes between different conditions [8]. It is independent of changes in the gene’s total expression, so it is not captured by gene-level differential expression analysis and requires transcript-level resolution to detect. 1, 4, 7, 20, 21, 25

droplet-based protocols (droplet data) High-throughput single-cell protocols that encapsulate each cell together with a barcoded bead inside a picolitre droplet, so that tens of thousands of cells are processed in one run. They are the dominant scRNA-seq format, and their capture chemistry reads only the tail of each transcript.

3

FASTQ The plain-text de facto standard output of Illumina sequencers, named by analogy to FASTA with a Q for quality. Files are almost always gzip-compressed, taking the .fastq.gz extension. 10, 23

full-length sequencing A sequencing approach that reads an RNA molecule from one end to the other, rather than only the fragment nearest one of its ends. 2

gene A functional unit of the genome that specifies one or more RNA molecules. 1, 23

genome The complete set of genetic material of an organism, written in DNA and inherited from one generation to the next. It is organised into genes, the functional units whose expression is measured. 1

Global Differentiation Index (GDI) A cluster-specificity measure used to identify marker genes and validate clustering quality. 19

isoform A distinct mature transcript variant produced from the same gene by alternative splicing. 1, 3, 4

isoform switching A change between conditions in the relative usage of a gene’s isoforms, such that the fraction of the gene’s expression contributed by one isoform increases while that of another decreases. It reflects altered alternative splicing or transcription and is quantified by the change in each isoform’s fraction of total gene expression [9]. 4, 20

messenger RNA (mRNA) The protein-coding subtype of RNA that carries a gene’s sequence to the ribosome. 1

mutual exclusivity Two features that rarely appear together in the same cell (a negative association), as opposed to co-expression. 20, 29, 35

orthogonal validation Independent confirmation of a computational result by a different experimental technique. 35

PCR Polymerase chain reaction: the amplification step that copies each captured cDNA fragment many times over, so that the few molecules recovered from a cell are

present in enough quantity for the sequencer to read [10]. Because copies outnumber original molecules, the UMI tag is what allows them to be collapsed back to a single count. 2, 10, 24

protein A molecule built from a chain of amino acids that performs the work of the cell, including catalysing reactions, transmitting signals, and forming structural components. 1

pseudo-bulk aggregation The practice of summing read counts across groups of single cells to simulate bulk RNA-seq depth, which recovers statistical power at the expense of single-cell heterogeneity. 4

read A short nucleotide sequence produced by the sequencing machine, representing the raw data of RNA-seq. 2

RNA (ribonucleic acid) The broad class of nucleic acid molecules transcribed from DNA. It comprises messenger RNA (mRNA), which carries protein-coding instructions, and non-coding RNAs that perform structural or regulatory roles. 1

RNA sequencing (RNA-seq) A technique that determines the sequence and abundance of RNA molecules in a sample; it is the basis of the single-cell protocols discussed in this thesis. 2, 41

single-cell RNA sequencing (scRNA-seq) RNA sequencing performed on thousands of individual cells separately, exposing per-cell expression. 2, 7, 8, 23

sparsity The very high proportion of zero counts in single-cell RNA-seq data. 3, 8, 9

splice junction The boundary created where two coding blocks (exons) are joined after the intervening non-coding block (intron) has been removed. Reads spanning a splice junction identify which blocks were joined together. 1

tag bias A bias in which each molecule is read only from one of its ends, so that coverage concentrates near that end while the interior of the transcript is never observed. It follows from a capture chemistry anchored at a single end, as in the tail-primed droplet protocols. 2, 3, 4, 9

technical dropout A zero count produced by a technical limitation of the assay—failed capture, reverse transcription, or sequencing—rather than by the true absence of the transcript. It is also described as a non-biological zero. 9

transcript An RNA molecule copied from a gene. 1, 4

Unique Molecular Identifier (UMI) A short random sequence attached to each captured RNA molecule before PCR amplification, so that amplified copies of one molecule are recognised and counted only once. 2, 10, 15, 24, 25

unspliced/spliced transcript The first RNA copy of a gene, still containing its non-coding blocks (introns); it is also called the primary transcript or pre-mRNA. Its counterpart is the spliced transcript, the same molecule once those blocks have been removed, which is the mature form translated into protein. 1, 27

Chapter 1

Introduction

Biological Context for Computer Scientists

Proteins are the basic building blocks of a cell and its DNA *genome* encodes the information needed to build them. Every cell of a multicellular organism carries essentially the same DNA, yet cells develop into distinct types. The difference cannot come from the DNA itself. Between the reading of the genome and the building of proteins, therefore, lies one or more layers of regulation that determine which proteins a cell produces, and those proteins in turn give the cell its characteristics. Studying these layers is central to explaining how cells function and what shapes their behaviour.

Gene expression proceeds in three stages. A *gene* is first copied into a long molecule of *RNA (ribonucleic acid)*, called *unspliced transcript*. The second stage, *alternative splicing*, removes the non-coding blocks and joins the coding ones at the resulting *splice junctions*, yielding a shorter spliced *transcript*. The final, spliced version is *messenger RNA (mRNA)*, which the ribosome translates¹ into protein in the third and final phase.

The shortening process is not a one-to-one mapping. To expand the decoding capacity of a gene, cells can produce several different spliced versions of the same *unspliced transcript*; these versions are called *isoforms*. Because different isoforms can translate into proteins with entirely different, or even conflicting, roles, a cell can fundamentally alter its behaviour simply by switching which variant it expresses—the MET Δ ₁₄ isoform, for instance, acts as a switch that drives invasive growth in cancer [11]. The core goal of *Differential Transcript Usage (DTU)* analysis is therefore to measure changes in the relative expression proportions of specific transcript isoforms from the same gene across different biological conditions, developmental stages, or cell types, in order to expose regulatory events that gene-level measurement renders invisible.

¹ **Translation:** The process by which the ribosome reads an mRNA molecule and assembles the corresponding chain of amino acids into a protein.

A Brief History of Single-Cell Sequencing

For most of its history, *RNA sequencing (RNA-seq)* was performed on bulk samples. RNA is extracted from a population of cells, pooled into a single reaction, and sequenced as one library², so the resulting profile describes the sample as a whole rather than any cell within it. Sequencing of this kind, termed *bulk RNA-seq*, is statistically comfortable: pooling thousands to millions of cells yields deep coverage and stable abundance estimates for every gene. The averaging, however, is also its principal limitation. A gene's measured abundance is the mean over all cells present, so cellular heterogeneity³ disappears before the data ever reach an analyst. A rare cell type constituting two percent of a tissue contributes two percent of that tissue's *reads* for each of its markers⁴ and is easily buried beneath the dominant population; two subpopulations that shift in opposite directions—one gaining a transcript while the other loses it—cancel out and appear unchanged.

Single-cell RNA sequencing (scRNA-seq) replaced that aggregate view with per-cell profiling, running the same measurement independently on thousands of cells. The strategy rests on compartmentalisation. Each cell is physically isolated, its polyadenylated RNA captured and reverse-transcribed separately, and every resulting molecule marked with two labels. A *cell barcode* identifies the cell of origin, so reads can be regrouped into per-cell profiles after sequencing; a *Unique Molecular Identifier (UMI)* tags the individual molecule, so *PCR (polymerase chain reaction)* duplicates are collapsed rather than miscounted as expression. A single pooled sequencing run is thereby converted into thousands of independent expression measurements, each far shallower than a bulk sample.

The first *scRNA-seq*, produced from 2009 onwards, came from plate-based protocols, which consume one cell per reaction [12]. The chemistry of that era anchored its reads at one end of the molecule, and base accuracy declines along a read as it lengthens.⁵ Reliable information therefore concentrated near that end, a limitation that deeper sequencing could not remedy—data of this kind exhibit *tag bias* [13].

In 2012, the invention of *full-length sequencing* removed that bias [14, 15]. A whole molecule could now be copied and read from end to end, preserving the internal structure that distinguishes one transcript variant from another. While this was a substantial gain in coverage, the technology was costly and the plate format kept a study at hundreds to a few thousand cells.

² **Sequencing library:** The set of DNA molecules loaded into the sequencer for one run. Here it holds the sample's RNA, copied into DNA and tagged so that the instrument can read each fragment.

³ **Cellular heterogeneity:** The coexistence of distinct cell types and states within a single organism.

⁴ **Marker genes:** A gene or transcript whose expression specifically identifies a cell type, state, or condition.

⁵ **Base:** One of the four nucleotides that form a DNA or RNA strand: adenine (A), cytosine (C), guanine (G), and thymine (T). A read is a sequence of consecutive bases, and the confidence with which each base is identified declines along a read as it lengthens.

A new protocol, droplet microfluidics, suddenly revolutionised the scale in 2015 [16]: each droplet encapsulates one cell together with a bead⁶ carrying barcoded primers⁷, so tens of thousands of cells are profiled in one run at a small fraction of the cost per cell. The gain came at a price. Droplet capture is anchored at the transcript tail; this forces the use of the old chemistry, and the tag bias returned.

Droplet protocols dominate single-cell work today, since they profile far more cells per experiment at a far lower cost per cell than plate-based chemistry allows [2]. What tag bias removes is therefore the central limitation of the format rather than an incidental one, and recovering it is an active line of work. Long-read platforms read whole transcripts at single-cell resolution, and short-read datasets are mined for whatever isoform signal their terminal fragments still carry [6].

The Computational Challenge

Two obstacles stand between droplet data and transcript-level results.

Extreme sparsity. *Count matrices* from droplet experiments are mostly zeros, with zero entries routinely forming 70% to 90% of all values. Some of those zeros are biological, recording a transcript that the cell never produced, while others are technical, recording a molecule that was present but never captured, and both look identical in the matrix.

Ambiguous reads. Only the terminal 3' end of each molecule is read, so where several *isoforms* share that end they look identical in the data and their proportions cannot be measured directly.

Deeper sequencing reduces the first obstacle and cannot touch the second: extra reads fill in more of the matrix, but no quantity of terminal fragments recovers the interior of a transcript. A method that works on droplet data must therefore extract signal from a lightly observed count matrix while accepting that most of each molecule was never seen.

A Note on What Remains Recoverable

Gene-level information survives the *tag bias* constraint. A gene registers when a read is confidently attributed to it, and in most cases a read that cannot be traced to a single transcript still maps to isoforms of the same gene. A gene matrix can therefore be built from the same reads, and it remains informative for *differential gene expression (DGE)* analysis. A consequential property is that this matrix is denser than the transcript matrix.

⁶ **Bead:** A small gel sphere loaded with the barcoded primers. One bead is co-encapsulated with each cell, supplying every primer its droplet needs.

⁷ **Barcoded primers:** Primers that carry a short identifying sequence, so that every molecule captured in a given droplet can be traced back to its cell after sequencing.

Current Approaches and Limitations

Three families of methods address transcript-level regulation, and each stops short of detecting DTU across a full droplet dataset.

Bulk algorithms. Established *bulk RNA-seq* tools such as rMATS [17], SUPPA2 [18], DRIMSeq [19] and DEXSeq [20] model continuous isoform proportions from deep, junction-spanning coverage. Applied to sparse droplet matrices, those assumptions break down and the resulting estimates are unstable.

Aggregation to pseudo-bulk. One route around the coverage problem is *pseudo-bulk aggregation*, which sums the cells of a cluster or condition into a single profile to recover the depth that bulk tools expect [3, 21]. The cost is the cell-to-cell variation the assay was built to measure: a switch confined to a small subpopulation is diluted by the cells around it.

Targeted 3'-biased sequencing tools. Sierra [4] analyses droplet data directly, but only for a narrow class of events at transcript termini. It offers no genome-wide framework for multi-*isoform switching* from raw zero-heavy counts.

What is missing is a framework that works on sparse droplet matrices at single-cell resolution, without collapsing cells together and without full-length reads.

Thesis Objective

The question this thesis pursues originates in COTAN (Co-expression Tables Analysis) [1], a framework for gene-level *count matrices* that focuses its statistical model on the zeros of the matrix rather than on its positive counts. It asks whether two genes are detected together less often than chance would predict, using presence/absence contingency tables, so no artificial imputation⁸ or variance-stabilising transformation⁹ is needed, and it performs well on the sparse matrices typical of single-cell data. This thesis asks whether the same statistics hold when applied one level lower, to *transcript* matrices whose counts are sparser still: whether they retain enough *isoform* information to detect *DTU* without fixing the *tag bias* problem. The claim is deliberately modest — isoforms that share their terminal region remain indistinguishable — and the aim is to map where sparsity-robust statistics recover differential transcript usage anyway.

The work proceeded from raw reads to a reported result in four stages, each producing a concrete artefact: a count-matrix pipeline, an analysis library, a reproducible driver layer, and a set of reported results.

⁸ **Imputation:** Replacing zero counts with values predicted from neighbouring cells, at the risk of inventing spurious correlations.

⁹ **Variance-stabilising transformation:** A transformation applied to counts, such as a logarithm, so that their variance no longer depends on their mean. It alters the scale of the data, turning an undetected molecule into a small positive value rather than a zero.

From reads to transcript matrices. A Nextflow pipeline was built that turns raw droplet reads into a filtered transcript-level count matrix—the artefact the whole study rests on, since DTU analysis requires transcript resolution that gene-level matrices do not carry.

An analysis library. The analysis was consolidated into a single R package, *cotanisoform*, which extends COTAN’s gene-level framework to transcript resolution. It keeps COTAN’s statistics—co-expression, p-values, GDI, differential expression—but adds the steps a transcript-level analysis needs: Seurat clean-up, COTAN object construction and annotation, clustering, and DTU candidate extraction. It is where COTAN’s gene-framed statistics become a transcript-level detector. The package also ships an improved logging layer that records progress and failures from both Seurat and COTAN in one consistent trace, so a long run reports its own state instead of failing silently.

The analysis drivers. A configuration-driven layer reproduces a full per-dataset analysis from a single declarative file, joining the pipeline and the library so the whole study runs end to end reproducibly.

From candidates to reported results. The workflow was applied to two datasets—a homogeneous human cell-line mixture and complex mouse cortical tissue—and the candidates are cross-referenced against published isoform biology and gene-level clustering baselines. The analysis was carried out first-hand and not reviewed by a domain expert; the findings are a reproducible first report rather than established biology, and independent verification is expected before they are built upon.

Chapter 2

Theory and Frameworks

This chapter provides the theoretical foundation for understanding the methods in Chapter 3. It covers *scRNA-seq* preprocessing concepts, the COTAN framework for sparse data [1], and *DTU* fundamentals.

Notation

Mathematical notation is used in this chapter and the two that follow only where it clarifies an argument. The table below collects the symbols that recur across them; symbols confined to a single local derivation are introduced where they appear.

Symbol	Meaning
X	Expression count matrix produced by the preprocessing pipeline, with features (transcripts or genes, depending on the pipeline output level) in rows and cells in columns
x_{fc}	Observed integer count of feature f in cell c
f, c	Feature and cell indices
N	Number of cells
R	Total number of sequencing reads
L	Sequenced read length
M	Mean number of reads per cell
k	Length of the k -mers used to index the reference
$ \Sigma $	Alphabet size (4 for DNA)
T	Number of annotated transcripts in the reference
T_{obs}	Distinct transcripts detected in a single cell
L_{avg}	Mean transcript length
θ_t	Relative abundance of transcript t
ℓ_t	Effective length of transcript t
C_r	Equivalence class of transcripts compatible with read r
$\hat{z}_{rt}$	Fractional assignment of read r to transcript t

Symbol	Meaning
<i>Statistics and detection</i>	
a_c, b_c	Binary presence indicators for a pair of features in cell c
O_{ij}, E_{ij}	Observed and expected counts in a 2×2 contingency table
Row $_i$, Col $_j$	Marginal totals of that table
χ^2	Pearson chi-squared statistic
$C_{f,f'}$	Co-expression coefficient between features f and f'
$p_{f,f'}$	Independence p -value for that pair
$D_{f,S}$	Differential expression (DEA) coefficient of feature f in cluster S
S	Cell cluster
τ	Contrast threshold on the DEA difference within a gene
gdi_i	Cluster-specificity index of feature i

Preprocessing Pipeline

This section follows a single *scRNA-seq* experiment from its data-generating mechanism to the *count matrix* consumed by every later analysis. Its scope is the route this thesis actually takes—droplet chemistry, pseudo-alignment, expectation-maximization—with the alignment alternative noted only where it changes the outcome. Plate-based and alignment pipelines differ in the details of each phase, not in the phase structure. It treats that path as a four-phase pipeline: ingestion and quality control, mapping, probabilistic resolution of ambiguous reads, and deduplication into a matrix. Each phase is mathematically and biologically constrained by the one before it, so the physical hardware chosen in the initial step dictates the shape, *sparsity* and content of the final analytical matrix.

Phase 0: Choosing the Protocol

The design of a *scRNA-seq* analysis pipeline is fundamentally constrained by the physical protocol chosen during data generation. Rather than representing a computational phase, this hardware choice establishes the data’s absolute boundaries, dictating the scale, resolution, and inherent *sparsity* of the resulting *count matrix*. This thesis follows the droplet route, whose three relevant consequences are the following. Capture is anchored at the 3’ terminus¹ of each transcript, so the internal sequence structure required to resolve complex splicing events is never observed. A finite sequencing capacity spread over a mass of cells reduces per-cell depth to a few thousand molecules. Empty droplets and multi-cell multiplets² must be filtered out downstream.

Table 2.2 sets the two chemistries side by side.

¹ **3’ (three-prime) terminus** is the end of a transcript at which the poly-A tail is added; capture chemistry anchored here reads only the terminal fragment, not the transcript body.

² **Multiplet:** A droplet that receives more than one cell, yielding a mixed profile that quality control must remove; also called a doublet when exactly two cells are involved.

Table 2.2: The two dominant single-cell library chemistries. Per-cell depth and cost trade off against throughput; total reads per day are set by the sequencer, not the chemistry. All figures describe a single sequencing run. Values are order-of-magnitude, after Ding et al. [22].

	Plate (SMART-seq2)	Droplet (10x)
Library preparation	~ 2 days for 96 cells	~ 1 day for 10^4 – 10^5 cells
Cells per experiment	10^2 – 10^3	10^4 – 10^5
Reads per cell	$\sim 10^6$	$\sim 10^4$
Detected transcripts per cell	$\sim 10^4$ – 10^5	$\sim 10^3$
Reads per day	$\sim 10^8$ – 10^9	$\sim 10^8$ – 10^9
Coverage	full transcript	3' terminus only
Reagent cost per cell	≈ 10 – $100\times$ higher	baseline

The defining computational consequence of the droplet architecture is extreme matrix *sparsity*, which follows directly from comparing the size of the reference with what a single cell yields. Let T denote the number of annotated isoforms in a standard mammalian reference and T_{obs} the distinct transcripts successfully recovered per single cell. While the reference dictates $T \sim 10^5$, a typical droplet-sequenced cell yields only $T_{obs} \sim 10^3$. In the raw alignment space, this disparity produces an initial matrix zero rate of roughly 99%. Standard preprocessing pipelines subsequently filter out features that are entirely unexpressed or ultra-rare in the sampled tissue, condensing the effective feature space down to $T_{filtered} \approx 1.2 \times 10^4$ – 2.3×10^4 viable transcripts. Within this practical, filtered domain, the observed zero rate consistently settles between 70% and 90%.

Some of these zeros are biological zeros³, recording a transcript the cell never produced; the rest are *technical dropouts*, recording a molecule that was present but never captured. Because the cumulative chemical efficiency of droplet capture, reverse transcription and sequencing succeeds on only 10% to 20% of the molecules present [23], the second kind dominates, and the two conflate into the 70% to 90% zero rate reported above. They are numerically identical once recorded, so no count-based test can separate them. The 3' *tag bias* deepens this sparsity further by collapsing distinct isoforms that share terminal exons into a single observational unit, and it is that overwhelmingly zero-heavy structure which dictates the need for a specialised statistical model in every downstream analysis.

The economics explain the adoption: a plate experiment consumes one well per cell at one to five dollars, while droplet encapsulation profiles a cell for a few cents, so a fixed budget buys two to three orders of magnitude more cells [2]. That trade is made knowingly. For the dominant questions—cell-type composition, marker discovery, differential gene expression—the payoff is statistical, because a per-gene test averages over thousands of cells and tolerates a sparsity that remains severe within any single one. Coverage loss is accepted wherever the question is gene-level, and resurfaces as a limitation only where it is isoform usage.

³ **Biological zero:** A zero count reflecting genuine absence of a transcript in a cell, as opposed to a technical dropout.

Phase 1: Data Ingestion and Quality Control

The first computational challenge in a single-cell pipeline is purely one of scale and structure. Sequencers output raw data in *FASTQ* format, representing each read across four distinct lines: a coordinate identifier⁴, the nucleotide sequence⁵, a separator⁶, and the corresponding Phred-scale quality scores⁷. For droplet-based chemistries such as the 10x Genomics platform, this raw sequence carries critical multiplexing metadata that must be carefully isolated before any biological analysis begins. Specifically, the sequencer must extract a *cell barcode* that maps the read back to its discrete cell of origin, and a *UMI*—a random tag attached prior to *PCR* amplification used to track individual captured mRNA molecules.

Before the reads can be quantified, they must be cleaned. Tools such as *cutadapt*⁸ or *fastp*⁹ trim synthetic adapters, discard low-quality bases, and extract the barcode and UMI payloads, while heuristic filtering separates valid cell barcodes from empty droplets and multi-cell collisions (multiplets). The step is optional, and is skipped under two conditions. When a vendor pipeline such as Cell Ranger has already trimmed adapters and corrected barcodes and UMIs, repeating the pass costs a full read of the *FASTQ* file without altering a base. When the mapper is read-structure-aware and UMI-aware, as *alevin-fry* and *Salmon* are, residual adapters fall outside the indexed k -mers and are never matched, so correction is left to the quantifier. Skipping is safe only under those guarantees: on untrimmed reads handled by a mapper that does not correct UMIs, adapters mis-assign reads and sequencing errors inflate the number of distinct UMIs. For plate-based chemistry no barcode or UMI exists to extract and adapter content is low, so trimming reduces to a quality filter.

Because quality control touches every single nucleotide exactly once, it acts as a highly parallelisable streaming operation strictly bound by the total read count, R , and the read length, L . The computational time complexity is entirely linear, scaling as:

$$\mathcal{O}(R \cdot L)$$

The physical footprint, by contrast, is large. One uncompressed *FASTQ* record—a 100 bp read together with its identifier and quality line—occupies roughly 250 bytes, so a droplet experiment of $R \sim 10^9$ reads, the order of magnitude considered here, holds $\sim 2.5 \times 10^{11}$ bytes, about 250 GB before compression. Modern multi-core trimmers sustain throughputs on the order of 10^8 reads per minute, so this I/O-bound filtering resolves in a few hours on a single workstation.

⁴ **Read identifier:** A line beginning with @ that names the read and records the instrument, flow cell and coordinates it came from.

⁵ **Nucleotide sequence:** The read itself, a string of A, C, G and T bases in the order the sequencer read them.

⁶ **Separator:** A line beginning with + that marks the end of the sequence and may repeat the identifier.

⁷ **Per-base quality score:** One character per sequenced base, encoding the probability that the base was called incorrectly on the Phred scale.

⁸ **cutadapt:** A read-trimming tool that removes adapter sequences and low-quality bases.

⁹ **fastp:** A fast all-in-one read preprocessor that performs adapter trimming, quality filtering and reporting in a single pass.

Phase 2: Mapping Algorithms (Alignment vs. Pseudo-Alignment)

Quantifying transcript expression is fundamentally a large-scale assignment problem: mapping short sequencing reads back to their reference transcripts of origin. The computational workload is strictly governed by the total read count, which scales as the product of the cell population and the sequencing depth. This relationship is defined mathematically as:

$$R = N \cdot M$$

where R is the total read count, N is the number of cells, and M is the average number of reads per cell. Both bounds below are derived with R symbolic and instantiated only once, here. The two datasets analysed in this thesis total roughly 10^8 – 10^9 reads, so $R \sim 10^9$ is the scale used throughout; pooled multi-sample atlases reach 10^{10} – 10^{11} reads and scale every cost in this section by the same factor. The search space and query geometry impose additional dimensional constraints, defined by the mammalian reference transcriptome and the physical sequencer output respectively:

$$T \sim 10^5, \quad L_{avg} \approx 1500 \text{ bp}, \quad L \approx 90\text{--}150 \text{ bp}$$

where T denotes the total number of annotated isoforms (including unspliced pre-mRNA), L_{avg} the average transcript length in base pairs, and L the sequenced read length in base pairs. Read length is strictly a hardware constraint; in a 3' assay, Read 1 isolates the cell barcode and UMI, leaving solely Read 2 to capture the actual transcript sequence.

Traditional alignment algorithms¹⁰ (e.g., BWA [24], Bowtie2 [25]) evaluate this mapping base-by-base to derive exact genomic coordinates using a Burrows-Wheeler Transform¹¹ index. Building this index requires a one-time computational investment proportional to the entire transcriptome size, scaling as $\mathcal{O}(T \cdot L_{avg})$. However, during the query phase, each of the R reads is localised through a search algorithm whose time cost grows logarithmically with the size of the reference T . Factoring in the algorithm's tolerance for sequencing errors, insertions, and deletions, this produces a total execution time of:

$$\mathcal{O}(T \cdot L_{avg}) + \mathcal{O}(R \cdot L \log T)$$

Searching the index does not scan the reference: backward search extends the matched suffix one base at a time, at constant cost per base. The logarithmic term above instead reflects the branch-and-bound over candidate loci that error tolerance requires, since a reference holding more transcripts offers more positions to test. That search penalty is what makes exhaustive alignment computationally prohibitive across the billions of reads generated in a single-cell experiment.

Modern droplet-based transcriptomics circumvents this bottleneck by utilising pseudo-alignment algorithms, such as Salmon, alevin-fry, or kallisto [26, 27, 5]. This approach rests on the critical mathematical observation that downstream quantification requires

¹⁰ **Read alignment:** Base-by-base mapping of a read to the reference, giving exact coordinates at high computational cost.

¹¹ **Burrows-Wheeler Transform (BWT):** A reversible text transform that allows fast substring search in a compressed reference index.

only the identity of the transcript of origin, not the exact base-level coordinates of the read. Consequently, pseudo-alignment replaces logarithmic sequence searches entirely with constant-time hash table lookups. Each read is fragmented into its overlapping k -mers¹², and every k -mer is queried in a pre-computed table recording which transcripts contain it. Intersecting these sets rapidly yields an equivalence class¹³ of candidate transcripts compatible with the read [5].

The procedure executes in two distinct computational phases. First, during index construction, the reference transcriptome is decomposed into an index—typically a colored de Bruijn graph¹⁴ mapping each k -mer to the specific transcripts containing it. This incurs a time cost identical to traditional alignment, but the index is stored sparsely: only the k -mers actually present in the reference are represented, giving a spatial footprint of $\mathcal{O}(d)$ for d distinct observed k -mers. The dense worst case, $\mathcal{O}(k \cdot |\Sigma|^k)$ over the alphabet of size $|\Sigma| = 4$, is $4^{31} \approx 4.6 \times 10^{18}$ entries and is never materialised. Like traditional alignment, this index is built exactly once before any read is processed.

Second, during the query phase, each read is fragmented into its $L - k + 1$ overlapping k -mers. The algorithm executes a hash table lookup for every single k -mer, retrieving and intersecting their corresponding transcript sets to produce the equivalence class. With a standard $k = 31$ and $L \approx 90$ –150 bp, this requires roughly 60 to 120 lookups per read. By shifting the core operation from a base-by-base sequence search to a hash retrieval, the per-read query cost collapses to $\mathcal{O}(L)$, completely decoupling it from the massive size of the reference transcriptome. The total execution time for pseudo-alignment therefore simplifies to:

$$\mathcal{O}(T \cdot L_{avg}) + \mathcal{O}(R \cdot L)$$

Instantiating the two query terms at $R \sim 10^9$ reads of length $L \sim 10^2$ bp against $T \sim 10^5$ transcripts, where $\log_2 T \approx 16.6$:

$$R \cdot L \log_2 T \approx 10^9 \cdot 10^2 \cdot 16.6 \approx 1.7 \times 10^{12}, \quad R \cdot L \approx 10^9 \cdot 10^2 = 1.0 \times 10^{11}$$

Two separate effects account for the gap. Pseudo-alignment removes the asymptotic $\log_2 T \approx 16.6$ factor, which is worth roughly one order of magnitude on its own. The remainder of the reported 50- to 100-fold speed-up comes from the constant factor: retrieving a hash-table entry costs far less than an inexact string match. The two effects multiply, and together they turn a query phase measured in days into one measured in hours.

¹² **k-mer:** A contiguous substring of length k ; reads are matched to transcripts by the k -mers they share.

¹³ **Equivalence class:** The set of transcripts that could have produced a read, since all contain the read's k -mers.

¹⁴ **Colored de Bruijn graph:** A graph of overlapping k -mers in which each k -mer is coloured by the transcripts containing it, used as the pseudo-alignment index.

Phase 3: Probabilistic Resolution via Expectation-Maximization

Pseudo-alignment introduces a distinct computational challenge in the form of multi-mapping reads¹⁵—reads whose k -mers match more than one transcript because those transcripts naturally share exonic sequences. A read falling entirely within a shared exon carries no intrinsic sequence information about which specific isoform it originated from, prompting the pseudo-aligner to return a set of valid candidates rather than a single definitive transcript. This ambiguity is profoundly exacerbated in $3'$ -*biased* protocols, where short reads capture exclusively the terminal region and frequently match every single isoform sharing that specific terminal exon [1].

Discarding these ambiguous reads is statistically unviable. Because they form the overwhelming majority of alignments in $3'$ -*biased* data, summarily dropping them would severely deplete the count matrix and artificially bias isoform proportions toward rare variants that happen to carry unique sequences.

Modern pseudo-aligners instead resolve this ambiguity by casting quantification as maximum-likelihood estimation over a generative model of the sequencing experiment [26]. Each fragment is assumed to be produced in two stages: a transcript t is first drawn with probability proportional to its relative abundance θ_t , and a start position is then drawn uniformly from the ℓ_t positions at which a fragment of the observed length can begin. The admissible count ℓ_t is the transcript's¹⁶. Crucially, the transcript of origin is latent: pseudo-alignment reports only the equivalence class C_r of transcripts compatible with read r [27]. Marginalising over the unobserved origin gives the probability of observing that class under a given abundance vector:

$$P(C_r | \theta) = \frac{\sum_{t \in C_r} \theta_t \ell_t}{\sum_{t'} \theta_{t'} \ell_{t'}},$$

and the quantity to be maximised is the product of these terms over all reads:

$$\hat{\theta} = \arg \max_{\theta} \prod_r P(C_r | \theta), \quad \sum_t \theta_t = 1.$$

Because the origin of each read is unobserved, this likelihood is optimised by the Expectation-Maximization (EM)¹⁷ algorithm, which alternates between two phases until convergence.

¹⁵ **Multi-mapping read:** A read compatible with more than one transcript because those transcripts share its sequence.

¹⁶ **Effective length:** The number of fragment start positions a transcript can generate; approximately its length minus the fragment length, plus one. It corrects for the fact that longer transcripts yield more fragments at equal abundance.

¹⁷ **Expectation-Maximization (EM):** An iterative algorithm that estimates transcript abundances and resolves ambiguous read assignments jointly until they converge.

In the E-step¹⁸, the current estimate is used to compute the posterior probability that read r originated from transcript t , given that it originated within the class C_r :

$$\hat{z}_{rt} = \frac{\theta_t \ell_t}{\sum_{t' \in C_r} \theta_{t'} \ell_{t'}} \quad \text{for } t \in C_r,$$

and $\hat{z}_{rt} = 0$ otherwise. The effective-length term ℓ_t is not optional: it encodes the fact that a long transcript contributes proportionally more fragments at equal abundance, so omitting it biases $\hat{\theta}$ toward short isoforms. The algorithm then proceeds to the M-step, where each abundance is re-estimated as the total posterior weight assigned to its transcript:

$$\theta_t \propto \sum_{r: t \in C_r} \hat{z}_{rt},$$

normalised so that the values of θ_t sum to exactly one across the transcriptome.

Two properties of this iteration matter for the interpretation of the resulting counts. First, each round is guaranteed not to decrease the log-likelihood, so the estimates converge monotonically to a *local* optimum rather than a guaranteed global one [28]. The likelihood surface is frequently flat or multimodal, so implementations initialise θ uniformly, stop when the relative change in log-likelihood falls below a tolerance, and may return different solutions for different starting points. Second, the model is closed-world: every read is assumed to originate from an annotated transcript, so an incomplete reference, unannotated isoforms or uncorrected fragment biases propagate directly into $\hat{\theta}$. The estimate is consequently a point estimate without uncertainty, and downstream methods that require variance—particularly the DTU tools discussed later—recover it by bootstrapping the quantification step [29].

Because the summation runs over equivalence classes rather than individual reads, many reads that share a compatible transcript set are handled by a single update. The computational cost therefore scales with the number of *distinct* classes, not with the number of reads, which is what makes transcriptome-wide EM tractable. Alevin-fry’s parsimonious variant exploits this further by collapsing transcripts that share an identical set of k -mers into a single group and running the EM updates over those groups, effectively degrading the problem to gene-level quantification where isoform resolution is impossible [5]. This is precisely the regime that *3'-biased* data tends to occupy, and it is worth stating why it cannot be repaired by more data.

¹⁸ **E-step / M-step:** The two alternating EM phases: distribute reads by the current abundance estimate (E), then re-estimate abundances (M).

Remark (identifiability under 3' capture). Let t and t' be two isoforms of the same gene whose reads fall only in equivalence classes containing both, as happens when the shared terminal region covers their effective lengths. The E-step above then assigns posterior weight through the ratio $\theta_t \ell_t / (\theta_t \ell_t + \theta_{t'} \ell_{t'})$ alone. With $\ell_t \approx \ell_{t'}$ that ratio reduces to $\theta_t / (\theta_t + \theta_{t'})$, so the likelihood depends on the two abundances only through their sum and is flat along $\theta_t + \theta_{t'} = \text{const}$: every split that preserves the total explains the observed reads equally well. Identifiability is therefore bounded by the divergence of the terminal region, not by the estimator, the reference or the number of reads, since sequencing the same fragment more deeply cannot separate two transcripts that fragment cannot distinguish. This is the formal reason the transcript-level analysis of this thesis can recover only those switches whose isoforms differ at or near their 3' ends.

Phase 4: Deduplication and Matrix Generation

Once each read has been assigned to a transcript, the remaining task is to count molecules rather than reads. The protocol attaches a *UMI*—a short random sequence—to each captured molecule before amplification, so every read carries both a *cell barcode* identifying its cell and a UMI identifying its molecule. Because PCR copies inherit both tags, the reads descending from one molecule are indistinguishable duplicates, and collapsing each such group recovers the molecule counts that sequencing had inflated.

The collapse must be performed on barcode, UMI *and* assigned transcript together. A UMI holds only a few random bases, so within one cell two unrelated molecules can share a UMI by chance; grouping on barcode and UMI alone would merge them, whereas including the transcript keeps them apart whenever they originate from different features.

The result is the expression count matrix X (Table 2.3), whose entry x_{fc} records the number of distinct molecules observed for feature f in cell c . Amplification no longer distorts the counts, so the differences remaining between cells reflect biology and capture efficiency rather than how often each molecule happened to be copied.

Table 2.3: The four phases of the preprocessing pipeline and the object each produces.

Phase	Operation	Output
0	Protocol choice	Library chemistry (plate or droplet)
1	Ingestion and quality control	Cleaned reads carrying a cell barcode and a UMI
2	Mapping (alignment or pseudo-alignment)	Reads assigned to candidate transcripts
3	Probabilistic resolution (EM, pseudo-alignment only)	Transcript-level abundance estimates
4	Deduplication and matrix assembly	Count matrix X

Figure 2.1 summarises the four phases and records what each one costs. The losses compound rather than compensate: the chemistry in Phase 0 decides which part of each molecule is readable at all, so every later phase works from that partial view. Phase 2

discards base coordinates, and Phase 3 leaves transcripts that share their terminal region statistically inseparable, which is the constraint quantified at the end of that phase. The figure is the reason Chapter 3 has to build a transcript-level matrix from raw reads: gene-level matrices lose the isoform distinction that Phase 3 could not recover in the first place.




Figure 2.1: The four phases of the preprocessing pipeline and the information each one removes. Phase 3 exists only on the pseudo-alignment route, where ambiguous reads must be distributed probabilistically; direct alignment resolves every read by its coordinates and skips it. The chemistry chosen in Phase 0 bounds every later stage, so the losses accumulate rather than cancel.

COTAN Framework

The transcript-level matrix X produced above acts as the natural input for the statistical framework presented here, yet it simultaneously introduces the primary analytical difficulty: the shallow per-cell depth heavily imposed by the droplet hardware leaves the vast majority of its entries at zero. The COTAN framework, introduced by Galfrè et al. [1], was specifically engineered to analyse this zero-heavy structure directly, and has since been logically extended to recover entire cell clusters purely from co-expression patterns [30]. The following subsections formalise this sparsity and detail the contingency-table machinery COTAN employs to successfully bypass traditional density-based statistics.

The Sparsity Problem

Single-cell count matrices exhibit extreme structural sparsity, with 70% to 90% of all entries returning zero [1]. As noted when the matrix was introduced, the biological and technical origins of those zeros are indistinguishable once recorded: a transcript that was truly absent and one that was merely missed leave the same mark [31]. That ambiguity, rather than the sparsity itself, is the fundamental statistical barrier for conventional tools designed for deep, dense *bulk RNA-seq* data.

Applying standard density-based statistical methods to such sparse matrices introduces severe, compounding computational artifacts. Log-normalisation, which strictly involves adding a pseudo-count before logarithmic scaling, heavily distorts low-abundance measurements and artificially compresses legitimate zeros into spurious scaling patterns [1]. Standard linear or rank-based association measures, such as Pearson or Spearman correlations, fail entirely in this domain, as they inevitably end up measuring the coincidence of technical dropouts rather than identifying true regulatory co-expression [1]. Furthermore, attempting to artificially densify the matrix through heuristic imputation—specifically replacing zeros with smoothed estimates derived from neighbouring cells—frequently introduces fabricated correlations and severely compromises sensitive downstream signals [1]. To fundamentally circumvent these pitfalls, COTAN explicitly abandons imputation and variance-stabilising transformations entirely, choosing instead to evaluate the zeros as a direct signal through strict binary discretisation.

Contingency Table Solution

COTAN cleanly bypasses the flawed continuous count amplitudes by evaluating pairwise feature dependencies using asymptotic independence testing applied strictly to binary presence and absence states [1]. For any given feature pair (f, f') across N cells, the continuous counts are logically reduced to discrete indicator variables for each cell c :

$$a_c = \mathbf{1}[x_{fc} > 0], \quad b_c = \mathbf{1}[x_{f'c} > 0]$$

Letting the discrete states be indexed by $i, j \in \{0, 1\}$, the framework constructs a comprehensive 2×2 *contingency table* where O_{ij} directly denotes the observed number of cells exhibiting state $a_c = i$ and $b_c = j$. The marginal totals for this specific table are defined as $\text{Row}_i = \sum_j O_{ij}$ and $\text{Col}_j = \sum_i O_{ij}$, mathematically constrained by the total cell count $\sum_i \text{Row}_i = \sum_j \text{Col}_j = N$. Under the standard null hypothesis of statistical independence, the expected cell count E_{ij} is calculated directly as:

$$E_{ij} = \frac{\text{Row}_i \cdot \text{Col}_j}{N}$$

Table 2.4 lays the four counts out explicitly. The two diagonal cells hold cells in which the features are detected together (O_{11}) or neither is detected (O_{00}); the off-diagonal cells hold cells defining the difference between them, in which exactly one of the two is present. Co-expression resides in the tendency of the count to fall on a diagonal rather

than on the anti-diagonal, and the expected value E_{ij} above is what that count would be under independence.

Table 2.4: The 2×2 contingency table built for a feature pair (f, f') across N cells. Entry O_{ij} counts cells with $a_c = i$ and $b_c = j$; the margins are the feature’s detection totals.

	f' detected ($b_c = 1$)	f' absent ($b_c = 0$)	Row total
f detected ($a_c = 1$)	O_{11}	O_{10}	Row ₁
f absent ($a_c = 0$)	O_{01}	O_{00}	Row ₀
Column total	Col ₁	Col ₀	N

Deviation from this independent expectation is evaluated using Pearson’s chi-squared statistic [1]:

$$\chi^2 = \sum_{i \in \{0,1\}} \sum_{j \in \{0,1\}} \frac{(O_{ij} - E_{ij})^2}{E_{ij}}$$

The result is reported not as a statistic but as a single coefficient. COTAN defines the co-expression coefficient $C_{f,f'}$ by rescaling the deviation so that it lies in $(-1, 1)$ and, under the null hypothesis, satisfies $N C_{f,f'}^2 \sim \chi^2(1)$ over the N cells of the table. The pairwise p -value, the quantity used for screening throughout this thesis, then follows directly:

$$p_{f,f'} = P(\chi^2(1) > N C_{f,f'}^2)$$

A positive coefficient means the two features are detected together more often than independence predicts; a negative one means less often, which is the mutual exclusivity that Chapter 3 treats as the signal of interest.

One modification separates $C_{f,f'}$ from the textbook Phi coefficient of a 2×2 table, and it lies in the expectation rather than in the statistic. Taking E_{ij} proportionally to the margins, as above, assumes that every cell has the same chance of expressing a given feature. Cells differ widely in library size, so the cells that capture less RNA accumulate zeros across all features at once, and a proportional expectation would read that shared under-detection as co-expression and inflate the coefficient. COTAN therefore derives the expected co-occurrence from a model of the zeros themselves: feature g in cell c is assigned a negative-binomial probability of a zero count with mean $\lambda_g \nu_c$ and shape α_g , where λ_g is the feature’s average expression, ν_c scales with the library size of cell c , and α_g is calibrated from the number of zeros actually observed for that feature [1]. Co-expression is thus measured against how often two features’ zeros were expected to coincide, not against their raw margins.

The chi-squared approximation behind $p_{f,f'}$ is asymptotic in the number of cells, not in the counts, and the distinction matters here. With N in the thousands the calibration is reliable for features detected in a reasonable fraction of cells, which covers most of a standard gene matrix. It degrades in the opposite regime, where a feature is rare enough that several expected cell counts E_{ij} fall below about five; Pearson’s statistic is then anti-conservative, and such a pair is best read as inconclusive rather than significant.

Filtering cannot remove this effect for the rarest transcripts, which is why candidates are ultimately ranked by effect size rather than by p -value alone.

A gene carrying k detected isoforms also contributes $\binom{k}{2}$ pairwise tests, and these are screened at a nominal $\alpha = 0.05$ with no correction applied. Since the tests are counted within genes rather than across the whole matrix, the adjustment a stricter procedure would impose is modest, and the effect-size threshold described in Chapter 3 provides the practical control. A permutation-calibrated null would be more principled and is noted as such in the discussion (Section 4.4).

Output Metrics

The framework reduces the machinery above to three output metrics [1], each of which is used later in this thesis. The first is the co-expression matrix C , whose entries are the coefficients $C_{f,f'}$ just defined; its negative values denote mutual exclusivity, and its positive values co-expression.

The second is differential expression analysis (DEA), which asks whether a single feature f is enriched in a given cell cluster S . It is the same construction applied to a different table: cells inside or outside S are crossed with feature f detected or not, and the resulting correlation coefficient $D_{f,S} \in (-1, 1)$ reports the association. A positive $D_{f,S}$ means f is detected more often inside S than outside it, a negative one less often. Because every DEA coefficient shares one bounded scale, values are comparable across features and across clusters, which is what allows the DTU contrast of Chapter 3 to be formed as the difference of two of them.

The third is the *Global Differentiation Index* (GDI), a per-feature cluster-specificity measure. For feature i it takes the mean of the smallest 5% of its pairwise p -values against all other features, $\bar{p}_i$, and applies $\text{gdi}_i = \ln(-\ln \bar{p}_i)$, which compresses the distribution of those means into an interpretable range [1]. Constitutively expressed features produce an approximately uniform lower tail and cluster around the null median $\ln(-\ln 0.025) \approx 1.305$, whereas features that vary across cell sub-populations accumulate very small p -values and rise above it. Values beyond roughly 1.5 therefore indicate features worth following up with DEA, and the GDI is used in Chapter 4 as the check that a clustering has recovered real population structure rather than noise.

Within transcript-level analysis the pairwise coefficient carries a sharper meaning than it does between genes. Observing a strong negative $C_{f,f'}$ between two isoforms of the same parent gene is the primary statistical indicator of alternative transcript usage, since it states that the two are detected in different cells more often than chance allows.

DTU Fundamentals

DTU vs Differential Gene Expression

Differential Transcript Usage (*DTU*) occurs distinctly when the relative expression proportions of multiple transcripts originating from the exact same parent gene change significantly between different biological conditions or cellular states. This concept differs fundamentally from Differential Gene Expression (*DGE*), which strictly measures overall changes in total, aggregated gene abundance without regard for internal composition. Biologically, DTU physically reflects shifts in *alternative splicing* decisions—a critical regulatory mechanism enabling cells to produce functionally distinct protein isoforms from an identical genomic sequence. Detecting these precise regulatory shifts inherently requires transcript-level resolution, as aggregating raw counts to the broader gene level mathematically obliterates the nuanced isoform switching signal entirely.

Why Bulk Methods Fail on scRNA-seq

The three compounding obstacles—shallow per-cell depth, the invisibility of internal splicing events under 3' bias, and the cell-to-cell variation that pseudo-bulk aggregation discards—are set out with the tools they affect in Section 1.4. Two structural points follow from them. Proportion-based DTU methods remain proportion models even when ported to single-cell data: the depth they require is still recovered by aggregation, or depends on junction-spanning coverage the droplet chemistry never provides [3, 21]. And as surveyed by Draper et al. [32], none of these contemporary approaches exploits the zero-heavy matrix pattern as a direct source of analytical signal. That gap provides the core motivation for adapting COTAN's sparsity-robust contingency framework to transcript-level analysis.

Mutual Exclusivity Signal

Within the adapted COTAN framework, negative statistical associations—formally defined as *mutual exclusivity*—between features serve as the foundational indicator of alternative usage patterns. When applied specifically to multiple isoforms of the exact same parent gene, this structural mutual exclusivity is biologically consistent with alternative splicing or *isoform switching* events, wherein different cellular states preferentially express entirely distinct transcript variants.

However, mutual exclusivity in isolation is mathematically insufficient to prove a functional switch. Two isoforms may yield a mutually exclusive signature simply because one variant is narrowly confined to a single specific cluster while the other is broadly ubiquitously expressed. This pattern of one-sided enrichment is biologically consistent with tissue-restricted isoform expression, but does not strictly reflect an active regulatory switch in splicing decisions. Identifying a genuine isoform switch requires explicitly

demonstrating bidirectional opposition: establishing at least one specific cell cluster in which the first isoform is heavily preferred, and at least one other distinct cluster in which that preference explicitly flips to its partner. This bidirectional contrast serves as the definitive conceptual filter mathematically separating a true *DTU* event from simple one-sided enrichment. While the present chapter successfully establishes mutual exclusivity as the core statistical signal worth testing, the precise mathematical metric that formalises this bidirectional contrast requirement is introduced and explicitly defined in Chapter 3.

Chapter 3

Materials and Methods

Data Acquisition and Dataset Selection

Two public droplet-based single-cell RNA-sequencing datasets form the empirical foundation of this study, each ingested as raw *FASTQ* reads together with its genome annotation.

Rather than searching the Gene Expression Omnibus (GEO)¹[33] ad hoc, candidate datasets were drawn from the curated collection at <https://seriph78.github.io/COTAN.Datasets.analysis/>. This resource offers an advantage that matters here: its datasets are pre-validated, since the same collections were used in COTAN’s original benchmarking papers, which confirms that they work with the framework and carry no major technical artifacts.

A structural constraint shapes what can be done with that collection: **no public *scRNA-seq* transcript-level count matrices exist**, at least within it. This reflects standard practice in the field rather than an oversight. Most pipelines, including CellRanger²[16] and the Seurat recommended workflow[34], collapse to *gene*-level counts by default; transcript quantification requires additional pseudo-alignment steps with Salmon[26] or alevin-fry[5] that many analysts skip; and even where transcript-level analysis is performed, the results are typically aggregated back to genes before publication. The consequence is that transcript matrices must be built from the *FASTQ* files directly.

That rebuild can be automated from scratch with `nf-core/scrnaseq` (v4.1.0)[35] (Section 3.2), an all-in-one, opinionated workflow that hides the wiring of each stage—indexing, alignment, quantification—behind a single invocation, lowering the entry barrier for an analyst new to the field.

¹ **Gene Expression Omnibus (GEO):** NCBI’s public database of published gene-expression datasets. It stores raw and processed data for re-analysis.

² **CellRanger / Seurat:** Standard pipelines (10x CellRanger, Seurat) that quantify and analyse *scRNA-seq* data, defaulting to gene-level counts.

That choice, however, narrows the field: `nf-core/scrnaseq` supports barcode-based protocols only, with the 10x Genomics[16] 3' chemistries, versions 2 and 3, as the supported family, so any dataset analysed here must come from that chemistry family. Of the COTAN datasets, only the Arigoni human cell-line mixture (GSE243665[36]) and the Ding mouse brain dataset (GSE132044[22]) satisfy both that chemistry requirement and the availability of published barcode lists. Those two datasets are therefore the ones analysed here.

Alignment, Quantification, and Feature Generation

Transforming raw sequencing data into transcript-level counts requires choosing a pseudo-alignment strategy and modifying its reference index to prevent premature aggregation.

Algorithm Selection: `alevin-fry` (simpleaf) vs. `kallisto`

Two pseudo-aligners dominate single-cell transcript quantification. Table 3.1 compares their primary characteristics.

Table 3.1: Comparison of scRNA-seq pseudo-alignment architectures.

Feature	<code>kallisto</code> – <code>bustools</code>	<code>alevin-fry</code> (simpleaf)
Multi-mapping handling	Uniform distribution	EM-based probabilistic assignment
UMI deduplication	Standard collapse	Directional error modelling
Execution speed	Fast (~10 min per sample)	Very fast (~5 min per sample)
<code>nf-core</code> integration	Available, non-standard	Native (v4.1.0)

The `simpleaf` framework[37] (v0.24.1 for quantification, v0.24.0 for the index build), which orchestrates `alevin-fry` for quantification and `piscem` for indexing, was selected for three reasons. Its UMI handling is superior, since single-cell data relies on UMIs to correct PCR amplification bias and the `alevin-fry` error model distinguishes true UMIs from sequencing errors more reliably. Reads compatible with multiple transcripts are assigned probabilistically via Expectation-Maximization (EM)[28] rather than split uniformly, which preserves the relative abundance signals that downstream analysis depends on. Finally, ³`nf-core/scrnaseq` (v4.1.0) supports `simpleaf` natively, with fewer configuration quirks than the `kallisto` alternatives.

The EM mode matters. A plain EM step spreads each ambiguous read across every transcript it is compatible with, so reads belonging to an abundant isoform leak counts onto the lowly-expressed isoforms that share its sequence, biasing the within-gene

³ **nf-core:** A community repository of standardised, peer-reviewed Nextflow analysis pipelines.

proportions that *DTU* detection reads out. *parsimony-em*⁴ instead groups reads that share an overlapping transcript set and first searches for the minimal set of transcripts that explains them, invoking EM only for the ties that remain genuinely unresolvable. Spurious low-abundance counts are suppressed and the estimated proportions are better identified. The mode returns fractional counts, which are rounded to integers before COTAN ingests the matrix.

Index Modification for Transcript-Level Resolution

Standard *alevin-fry*[5] output collapses to gene level by default via transcript-to-gene mapping⁵. To extract transcript counts, the index files were modified before alignment, following the identity-mapping route discussed in the *simpleaf* issue tracker⁶:

1. **Identity mapping in `t2g_3col.tsv`:** the standard mapping is replaced by a self-referential one, which prevents the final aggregation step that sums isoforms into genes.

```
# Original: ENSMUST00000193814  ENSMUSG00000067548  # parent gene
# Modified: ENSMUST00000193814  ENSMUST00000193814  # itself
```

2. **Removal of `gene_id_to_name.tsv`:** the *nf-core* pipeline uses this file to map Ensembl gene identifiers to gene names during annotation. With the identity mapping in place the matrix rows are Ensembl *transcript* identifiers, so the annotation step finds no gene IDs to match and halts. It halts after the transcript counts have already been written to disk, so those counts are kept and carried into the analysis.

This modification is **reversible**: transcript-level quantification is a finer partition of the gene-level signal, so summing the transcript rows by gene using the retained `t2g` and `gene_id_to_name` tables recovers a gene-level matrix directly, without rebuilding the index or re-aligning. The tools are not designed for this detour, and the collapsed matrix matches a native gene-level run except where the two resolutions assign ambiguous reads differently.

⁴ **parsimony-em:** A *UMI*-resolution mode that restricts each multi-mapping read to the minimal set of transcripts explaining it, applying EM only to the ties that remain.

⁵ **Transcript-to-gene (t2g) mapping:** A file pairing each transcript with its parent gene; replacing it with an identity mapping prevents gene-level collapse.

⁶ <https://github.com/COMBINE-lab/simpleaf/issues/115> (accessed 29 September 2026)

Reproducible Workflow Automation

To make this workflow reproducible across datasets, it was automated with **Nextflow**[38] (v26.04.4), a domain-specific language for computational pipelines that handles parallelization and checkpointing. The resulting wrapper calls `nf-core/scrnaseq` (v4.1.0), a community-maintained scRNA-seq analysis pipeline with built-in support for `simpleaf`. This choice provides error handling, containerized dependencies, and standard configuration across datasets.

The custom automation script:

- **Input:** a CSV samplesheet listing each sample and its SRA accessions⁷, and a per-dataset configuration naming the genome assembly and species, the Ensembl release (102), the 10x chemistry, and the published barcode list. The pipeline downloads the reference FASTA and GTF⁸ from Ensembl itself, so a matching assembly name replaces a manual reference download.
- **Automated steps:**
 1. Download the FASTQ files via `sra-tools`[39].
 2. Download the Ensembl reference and build the `piscem` index, in either a gene-level or a transcript-level mode.
 3. Apply the identity-mapping modification when transcript-level output is requested.
 4. Execute `nf-core/scrnaseq` (v4.1.0) with `simpleaf`.
 5. Produce the raw count matrix and its metadata.
- **Availability:** the complete pipeline, including configuration files and test cases, is available at <https://github.com/ldelitala/sc-isoform-cotan>.

The workflow ends at the count matrix. Quality filtering and the COTAN analysis are performed afterwards in the R analysis layer (R 4.5.3), keeping the two concerns separate.

Adapter trimming is omitted deliberately rather than by oversight. The 10x read structure places the barcode and UMI in a fixed header segment, and `simpleaf` resolves that geometry internally with a UMI-aware error model, discarding residual adapter k -mers because they match no indexed sequence. A separate `fastp` pass would therefore reread the entire FASTQ file for no change in the count matrix. Quality control at the barcode level still runs: `simpleaf` applies empty-droplet detection to the quantification output, and cells are further restricted to the high-quality barcode sets published with each dataset (Section 3.3).

⁷ **Sequence Read Archive (SRA):** NCBI's repository of raw sequencing data, accessed by identifiers such as SRR accessions.

⁸ **GTF (Gene Transfer Format):** An annotation file listing each gene's coordinates, exon structure, and transcript boundaries; used to build the transcriptome index.

Data Processing and Quality Control

Cell-Level Filtering

Cells were filtered using metadata from the original GEO publications. The Arigoni dataset (GSE243665[36]) and the Ding dataset (GSE132044[22]) both provide lists of cells that passed their quality control thresholds, and restricting the analysis to those known-good barcodes avoids a second QC pipeline that could introduce inconsistencies. Published studies typically filter on a minimum UMI count per cell, which removes empty droplets; a maximum mitochondrial read percentage⁹, which removes dying cells; and a gene detection threshold, since cells with too few detected genes may be doublets or otherwise low quality. Reusing their filtering ensures comparability and reduces methodological drift.

The unfiltered matrices carry 488,342 (Arigoni) and 121,082 (Ding) barcodes, most of them empty droplets; restricting to the barcodes the original studies report as passing their quality control leaves **29,606** human and **4,091** mouse cells.

Feature-Level Filtering

Feature filtering ran in the R analysis layer (R 4.5.3) with COTAN (v2.13.1, commit be93aa8) for both the spliced-feature restriction and the `clean()` routine, and Seurat (v5.5.1) as the count-object container. Transcript matrices contain many features that add noise without signal. *unspliced transcripts* are removed because COTAN’s framework[1] was designed for mature mRNA, whereas unspliced pre-mRNA has different statistical properties and can confound *co-expression* analysis; filtering retains spliced-only features and reduces matrix dimensionality. The remaining sparsity is trimmed by COTAN’s `clean()` routine, which drops every transcript expressed in fewer than 0.2% of cells and every cell detecting fewer than 0.3% of the retained features—a fractional cutoff far stricter than a fixed count, because isoform-level detection is so sparse that an absolute threshold would admit thousands of near-empty rows. A handful of fully-expressed transcripts are removed as well. After filtering, the human matrix holds 29,454 cells $\times$ 23,673 transcripts and the mouse matrix 4,077 cells $\times$ 12,848 transcripts (the residual 152 human and 14 mouse cells fall below the sparsity cutoff). Although the transcriptome lists more transcripts than genes, the cleaned transcript matrix ends up smaller than the gene matrix: the raw mouse matrix detects 149,443 transcripts, yet after filtering only 12,848 remain, against 28,692 genes in the authors’ GEO count set—isoform-level detection is so much sparser that most transcripts never pass the sparsity cutoff.

These choices prioritize quality over quantity. Smaller matrices mean faster computation and cleaner signal—important when running expensive contingency table calculations across all feature pairs.

⁹ **Mitochondrial read percentage:** The fraction of a cell’s reads coming from mitochondrial genes; abnormally high values flag dying or damaged cells.

Before the COTAN object is built, the Seurat assay’s counts are rounded to the nearest integer and explicit zeros are dropped, since COTAN ingests integer counts; the fractional parsimony-em assignments never enter the framework directly.

Population Partitioning and Validation

Differential transcript usage is a contrast between groups of cells: an isoform switch surfaces only when the cells carrying one variant can be set against the cells carrying the other. The partition of cells fixes what the detector is able to see, and the two datasets reach that partition along different routes.

Ground-Truth vs. Computational Partitioning

The Arigoni human cell-line dataset arrives already divided: its GEO deposit distributes one barcode list per constituent cell line and for the PBMCs, so those published labels serve directly as the groups the detector compares.

The Ding mouse cortex dataset arrives with no such labels. Its cells are grouped from the matrix itself with COTAN’s hierarchical uniform clustering¹⁰. Cell-to-cell dissimilarity is read from the co-expression patterns, clusters are split until each becomes homogeneous, and neighbours the statistics cannot separate are merged. Run on transcript counts, the procedure returns eleven uniform clusters and a small noisy remainder—the cells whose expression matches no group closely enough to claim them. Those eleven clusters are a hypothesis, not a verdict: an isoform matrix is thinner than a gene matrix, and nothing guarantees that the co-expression structure survives the thinning.

Cross-Resolution Validation

To give the transcript partition something to be measured against, the same cells were clustered a second time on gene counts: the official Ding gene matrix (4,091 cells $\times$ 28,692 genes), cleaned with COTAN’s standard `clean()` and clustered with parameters identical to the transcript run. Gene-level clustering is the established use of COTAN and, at four thousand cells, costs a fraction of the transcript run. It returns fifteen clusters where the transcript run found eleven.

A shared skeleton with local departures is the outcome the method hopes for. The two runs are not the same cells at two aggregation levels: the gene partition comes from the independent GEO gene matrix, which carries 28,692 features against the transcript matrix’s 12,848, so its fifteen clusters reflect a denser feature set rather than a finer resolution. The informative evidence is not the total counts but where the partitions

¹⁰ **Hierarchical uniform clustering:** COTAN’s procedure that recursively splits cells by co-expression dissimilarity, then merges statistically indistinguishable clusters.

agree and where they split or merge—near-identical grouping would leave the isoform resolution unjustified, and a systematic departure is where transcript clustering earns its keep.

The two partitions meet over the 4,077 cells common to both analyses. A confusion matrix¹¹ cross-tabulates their labels: a dominant diagonal would say that both methods recover the same structure, and the cells that fall off it expose where one partition splits a group the other keeps whole. To make the comparison depend on signal rather than on the partition, the detector is then run twice over the same transcript matrix, once with the transcript clusters and once with the gene-derived clusters standing in as ground truth, so that an event surviving both is a signal no choice of partition can manufacture, while an event unique to one belongs to the resolution that found it. Event overlap closes the reading: the genes and transcripts both runs call positive are the robust switches, the transcript-only calls are the subtle ones that need fine partitioning to surface, and the gene-only calls are the faint ones that lean on the statistical power of aggregation.

Neither partition can be declared the correct one, since the mouse dataset carries no reference labelling to check against. What the comparison delivers is not a winner but a description of what each partition can see, and of the biology that comes into view only when the cells are cut at isoform resolution.

The DTU Detection Algorithm

The detector bypasses pseudo-bulk aggregation and instead reuses COTAN’s gene-level machinery—the contingency table, the *co-expression* coefficient, and the differential-expression contrast—applied at transcript resolution, so that a mutually exclusive isoform pair is recognised by the same independence test that already copes with sparsity[1]. The biological justification for the two filters is given in Section 2.3.3. For every gene with $k \geq 2$ detected isoforms, the procedure is:

1. ***mutual exclusivity screening***: for every possible pair of isoforms (A, B) within the gene:
 - Compute the *co-expression* coefficient $C_{A,B}$ from the 2×2 *contingency table* of presence and absence across cells[1] (Section 2.2).
 - Retain pairs with a negative $C_{A,B}$ and a p -value¹² < 0.05 .
2. ***Contrast calculation***: for each significant pair, compute the differential enrichment vector

$$\Delta_{\text{cluster}} = D_{A,S} - D_{B,S}, \quad (3.1)$$

¹¹ **Confusion matrix**: A table of cluster-label agreement between two clustering runs, revealing systematic splitting and merging.

¹² **p-value**: The probability of observing an association this strong if the features were truly independent.

where $D_{f,S}$ is COTAN’s differential expression coefficient for feature f in cluster S (Section 2.2). Each term lies in $(-1, 1)$, so the contrast is bounded and comparable across genes and clusters.

3. **Bidirectional thresholding:** retain pairs that switch, which requires opposition in both directions:
 - At least one cluster with $\Delta > \tau$.
 - At least one different cluster with $\Delta < -\tau$.
 - ¹³ τ is calibrated per dataset in the section that follows.

The algorithm outputs a list of gene-isoform pairs with evidence for cluster-specific differential usage, ranked by statistical significance and effect size ($|\Delta|$).

Calibrating the Switch Threshold (τ)

The contrast threshold τ controls the balance between sensitivity and specificity: high values yield few confident calls but risk missing subtle switches; low values capture more events but increase false positives.

A tuning script was developed at the supervisor’s request (Prof. Silvia Galfré). It varies τ across a range such as 0.01 to 0.5, counts the detected DTU events at each threshold, plots the relationship between τ and the number of significant pairs, and suggests thresholds that yield biologically reasonable candidate counts, typically 20–100 events for validation feasibility.

The script is included in the GitHub repository alongside the main pipeline. Results from this tuning informed the choice of $\tau = 0.2$ for human data (well-separated cell lines, stronger signals) and $\tau = 0.05$ for mouse brain tissue (continuous differentiation trajectory, subtler transitions).

Detailed threshold optimization results appear in Chapter 4 alongside the DTU discovery findings.

¹³ τ (**contrast threshold**): The minimum absolute DEA contrast between two isoforms required to count a cluster toward a switch.

Chapter 4

Results and Discussion

This chapter presents the findings from applying transcript-level COTAN to scRNA-seq data. It begins with quality assessment (GDI clustering validation) and a comparison of gene-level and transcript-level structure, before presenting the main contribution: DTU detection results across the human and mouse datasets. Each section builds toward answering whether sparse 3'-biased data can support isoform-switch discovery using contingency-table statistics.

Dataset Quality Assessment via GDI

Before DTU calls can be interpreted, the clustering must be shown to capture biological structure rather than technical noise. The Global Differentiation Index (GDI), defined in Section 2.2, provides this validation.

Transcript-level Clustering Quality

Figure 4.1 shows the GDI plot for the mouse dataset clustered on transcript counts. Each point represents a cell, colored by its assigned cluster (11 uniform clusters total). The clear spatial separation indicates that transcript-level co-expression patterns distinguish distinct cellular populations.

Key observations from this visualization:

- Cluster boundaries are well-defined with minimal overlap, suggesting robust isoform-based signatures.
- High-GDI genes (marker candidates) show selective enrichment in specific clusters rather than uniform expression.
- The noisy remainder cells (not assigned to any cluster) form a tight central cloud; these were filtered out from downstream DTU analysis.

This confirms that transcript-level data contains sufficient signal for meaningful clustering, despite sparsity and 3'-bias constraints.




Figure 4.1: GDI plot of the mouse dataset clustered on transcript-level counts, coloured by the 11 uniform clusters.

Gene vs Transcript Structure Comparison

Figure 4.2 shows the same cells clustered using standard gene counts (15 clusters). The gene run reads the independent GEO gene matrix—28,692 features against the transcript matrix's 12,848—so its larger cluster count reflects the denser matrix, not a finer gene-level resolution.

These GDI analyses establish that both clustering approaches capture real biological signal, but with different sensitivity to cellular substructure. The main question is whether this structural difference translates into distinct DTU detection.

Clustering Comparison: Gene vs Transcript Level

To quantify how gene and transcript clustering diverge, a confusion matrix was constructed across 4,077 common cells present in both analyses.

Figure 4.3 shows the confusion matrix heatmap. Several patterns emerge:

- **Diagonal dominance:** Most cells receive consistent assignments across both methods, validating that neither approach produces random clusters.




Figure 4.2: GDI plot of the same cells clustered on gene-level counts, yielding 15 clusters.




Figure 4.3: Confusion matrix comparing transcript-level (11 clusters) and gene-level (15 clusters) assignments for the 4,077 common cells of the mouse dataset.

- **Systematic splitting:** Several transcript-level clusters correspond to sub-regions of single gene-level clusters. For example, transcript cluster 4 splits from the broader gene cluster 7, suggesting isoform variation within what appears homogeneous at gene resolution.
- **Selective merging:** Conversely, some adjacent transcript clusters merge in gene analysis, indicating that their differences rely on isoform-specific expression patterns rather than overall transcriptional activity.

The interpretation: across the two runs the counts differ (11 against 15) because the partitions come from different matrices—the transcript matrix is a sparser, feature-filtered view, the gene matrix a denser GEO count set—so the totals are not a resolution ranking. What the confusion matrix actually shows is agreement in the large and divergence in the particular: the dominant diagonal confirms both partitions recover the same broad populations, while the local splits and merges mark where isoform information changes the grouping. That is precisely the regime where DTU detection should add value—if isoform switches distinguish cell states that genes alone cannot separate, then transcript-aware methods capture biologically relevant signals.

However, this raises a critical question: whether the detected DTU events hold across clustering strategies, or depend entirely on how cells are partitioned. The robustness analysis below answers it.

Core Finding: DTU Detection Results

The algorithm from Section 3.5 was applied to both human and mouse datasets, filtering for mutually exclusive isoform pairs with bidirectional cluster preference (threshold τ as specified). The results show that contingency-table statistics can recover within-gene exclusivity patterns even in sparse 3'-biased data.

Human Dataset Findings (Arigoni)

The human cell line mixture provided ground-truth validation through known biology. Of the 4958 genes carrying at least two detected transcripts, 430 retain at least one significant mutually exclusive pair¹, yet after the bidirectional contrast threshold $\tau = 0.2$ only 21 candidates survive, across just 9 genes. That collapse from 430 to 9 is the algorithm working as designed: mutual exclusivity alone admits hundreds of pairs, but only a handful also flip their enrichment between clusters, and it is this second condition—a pair must be preferred in one cell type and disfavoured in another—that separates a genuine switch from a statistical curiosity.

¹ Counted by `extract_dtu_candidates()` and recorded in the run log (`results/arigoni/logs/`); the exported `dtu_candidates.csv` holds only the 21 survivors.

The surviving candidates cohere into a single story. TPM1, the strongest call, is a tropomyosin whose isoforms oppose each other across the cell lines: one isoform is enriched in PC9 and DV90 while the other dominates A549 and HCC78, a reciprocal preference that reads as cell-type-specific cytoskeletal tuning. Alongside it sit the ribosomal proteins RPL10 and RPS24, whose unexpected DTU signal points toward cell-type-specific translation regulation—consistent with the heterogeneity of ribosomal composition underlying specialised translation [40]. The two themes are biologically distinct but converge in method: in both cases a single gene carries two isoforms that the same cell types use in opposite directions, and it is that convergence, not any single gene, that validates the detector. The final list is small precisely because the algorithm demands both *mutual exclusivity* and a bidirectional contrast, and that conservative design is what lets nine genes carry the whole human result.

Mouse Dataset Findings (Ding)

For mouse brain tissue a lower threshold, $\tau = 0.05$, was applied, and the balance shifts: only 14 genes carry significant pairs, yet 46 candidates survive across 8 genes. The inversion is itself informative—fewer initial pairs but more survivors means mouse cortex switches are gradual rather than sharp, so a relaxed threshold admits many weak but real transitions that the strict human cutoff would have discarded.

The candidates trace an arc across the tissue. Meg3, a long non-coding RNA² with a documented role in neuronal differentiation, switches along the differentiation trajectory itself: isoform A dominates early progenitor clusters while isoform B is enriched in differentiated neurons—a transition that tracks a continuous gradient rather than a binary state. Malat1, another highly conserved lncRNA involved in splicing regulation, shows the same cluster-specific preference, consistent with the alternative 3' isoform usage documented for this locus [41]. The two lncRNAs are neighbours in function—both non-coding, both switching with differentiation—and together they anchor the mouse result in established biology. The arc then turns inward: Srrm2 encodes a core spliceosome³ component, and its own DTU suggests feedback regulation, alternative splicing of splicing factors being a known mechanism for tuning isoform landscapes across cell states. Finally, Ttc14, a novel candidate with no prior DTU literature, requires *orthogonal validation* to confirm its biological significance—and its presence is a reminder that the method also surfaces calls no published study has anticipated.

² **Long non-coding RNA (lncRNA):** A long RNA molecule that does not encode a protein but regulates gene expression, splicing, and development.

³ **Spliceosome:** The molecular machine that removes introns and joins exons during splicing.

Robustness Analysis: Shared vs Exclusive Events

DTU detection was re-run using gene-derived cluster labels on the same transcript matrix, and the resulting event lists were compared across both clustering strategies. Both partitionings returned the same total, 46 DTU pairs; their composition differed:

- **Shared events:** Majority of pairs detected by both methods, including *Meg3* and *Malat1*—these represent strong signals surviving regardless of partitioning scheme.
- **Transcript-exclusive (3 pairs):** *Srrm2* and *Ttc14* isoform switches appeared only with transcript-derived clusters. These likely reflect subtle transitions requiring fine-grained population resolution to detect.
- **Gene-exclusive (4 pairs):** *Snrnp70*, *Lrrc7*, *Zfp941*, and a specific *Meg3* pair detected only using gene-level clustering. These may represent weaker signals that benefit from aggregated statistical power.

The interpretation:

1. Strongest DTU signals (high effect size) survive both approaches—validating the algorithm’s core functionality.
2. Subtle switches depend on clustering resolution: transcript-level captures isoform-specific states; gene-level provides power for low-abundance transcripts.
3. Neither approach dominates universally. Researchers should consider their biological question when choosing—fine-grained substructure vs statistical robustness.

This robustness analysis confirms that while the strongest signals are cluster-agnostic, subtle isoform switches reveal themselves only under appropriate partitioning schemes. The choice of clustering strategy is not merely technical—it shapes which biology becomes visible.

Threshold Sensitivity: The Tau Sweep

The candidate count depends on the contrast threshold τ . To make the choice principled rather than arbitrary, the analysis re-runs the DTU extraction over a fine grid of τ values (0.005 to 0.3 at step 0.005) and records, at each threshold, the number of accepted events, the unique genes behind them, and how many survive the strictest threshold that still yields at least ten events (the robust core).

Two properties of the sweep matter for interpretation. First, the event count is monotone decreasing in τ —a sanity check that the extraction responds predictably to the threshold. Second, the fraction of accepted events belonging to the robust core (`frac_in_core`) rises toward one as τ approaches the core threshold; a plateau in this fraction marks the point where further relaxation adds only weak, non-core switches.

Figure 4.4 maps that search across a landscape of thresholds. The bulk of significant pairs descends like a long, smooth slope—from 2006 at $\tau = 0.005$ to 21 at $\tau = 0.2$ —a nearly featureless gradient that reflects the homogeneous cell-line mixture: every isoform pair carries some reciprocal contrast, so the terrain is a dense and continuous spectrum with no natural break. On such a slope any single threshold looks arbitrary, and this is exactly the trap the sweep is built to expose. As the eye climbs the landscape a highland finally appears: at $\tau = 0.25$ ten events, all belonging to the robust core, hold flat through $\tau = 0.26$ —a genuine plateau, ground stable enough to stand on rather than a point on the slope. The published threshold $\tau = 0.2$ reproduces the 21 reported candidates and sits just below that highland, so the human result is anchored at the edge of stable ground rather than lost somewhere on the tail. The mouse terrain is different: it descends in steps, 46 events at $\tau = 0.05$ holding at 41 and 36 across wide intervals, because heterogeneous tissue is dominated by a few high-contrast switches with a sparse tail—so its plateaus are broad ledges, not a single summit.




Figure 4.4: Tau sensitivity sweep for the human dataset: DTU events (blue) and robust-core events (orange) versus the minimum DEA contrast threshold τ . A stable plateau appears at $\tau = 0.25$ with a ten-event core; the published $\tau = 0.2$ (dashed) reproduces the 21 reported candidates.

The plateau provides an empirical anchor for τ : the reported thresholds (0.2 human, 0.05 mouse) were chosen within or adjacent to stable regions, so the final candidate lists are not an artifact of a single fragile threshold value. The sweep does not, however, replace a permutation-based null distribution; it remains an internal-consistency check on how the accepted set changes with the threshold.

Discussion

These findings are now set within their broader methodological and biological context.

Key Findings Interpretation

Two conclusions follow from these results. The top candidates (Meg3, Malat1, TPM1) align with documented biology—lncRNA switching in neurogenesis and ribosomal protein regulation—which indicates that the algorithm recovers genuine signal rather than technical artifacts. Convergence across a human cell line mixture and mouse brain tissue further indicates that the method holds under different biological contexts rather than a single favourable case study.

Limitations and Caveats

Despite encouraging results, several constraints temper interpretation:

- **3'-bias invisibility:** Alternative splicing events occurring in transcript mid-regions remain undetectable when only 3' ends are sequenced. The DTU candidates found here represent a subset—those where isoforms differ at their terminal exons or 3' ends.
- **EM rounding effects:** Fractional EM assignments were rounded to integers for COTAN compatibility (Section 3.3). Low-abundance isoforms with small fractional counts may be distorted by this discretization, potentially introducing false negatives or altered contrast values.
- **Threshold sensitivity:** The choice of τ lacks formal statistical grounding. The tau sweep (Section 4.3) anchors the threshold empirically at a stable plateau, but a permutation-based null distribution would still provide calibrated p -values rather than a consistency-based choice.
- **Lack of orthogonal validation:** Without long-read sequencing or targeted RT-PCR⁴ confirmation, false positive rates remain estimated from internal consistency (shared vs exclusive events) rather than ground truth. This is a limitation common to computational DTU studies but remains critical for biological claims.

⁴ **RT-PCR:** Reverse-transcription polymerase chain reaction: a lab technique that experimentally confirms whether a specific isoform is present, used as an orthogonal check on computational DTU calls.

Future Directions: Extending COTAN for Transcript Analysis

The framework established here opens several avenues for methodological extension:

1. **Integration with long-read data:** Combining short-read depth (thousands of cells) with PacBio/Nanopore⁵ isoform annotation would anchor DTU calls in complete transcript sequences.
This hybrid approach could validate computational predictions while scaling to population-level sample sizes.
2. **Trajectory-aware analysis:** Current clustering assumes discrete cell populations, but differentiation is continuous. Modeling isoform proportions along pseudotime⁶ trajectories (e.g., using COTAN’s framework adapted for regression rather than categorical comparison) could reveal dynamic switching patterns during transitions.
3. **Allele-specific expression extension:** COTAN’s contingency table approach naturally extends to phased variants. Detecting allele-biased isoform usage would connect regulatory genetics (eQTLs, sQTLs)⁷ to splicing outcomes in single cells—bridging population genomics with transcript regulation.
4. **Benchmarking against bulk DTU tools:** Systematic comparison with rMATS, SUPPA2, and DRIMSeq [29, 32] on matched pseudo-bulks would contextualize sensitivity and specificity. This benchmark could identify regimes where single-cell methods add unique value versus when aggregation suffices.

These extensions remain beyond the scope of this thesis, but they indicate that the framework admits further adaptation as new data types and biological questions emerge.

⁵ **Long-read sequencing (PacBio, Oxford Nanopore):** Sequencing technologies that produce full-length transcript reads, giving complete isoform annotation unavailable from short 3’-biased reads.

⁶ **Pseudotime:** A one-dimensional ordering of cells along a continuous developmental or differentiation trajectory, capturing gradual transitions rather than discrete populations.

⁷ **eQTL / sQTL:** Genomic variants that measurably affect gene expression (eQTL) or splicing/isoform usage (sQTL); linking them to single-cell isoform switches connects population genetics with regulation.

Chapter 5

Conclusions

This thesis established that COTAN’s sparsity-robust contingency-table framework can be extended to transcript-level analysis, recovering within-gene exclusivity patterns indicative of Differential Transcript Usage (DTU) from sparse 3’-biased single-cell *RNA-seq* data.

Summary of Findings

Applied to the two droplet datasets, the adapted framework returned detectable signal in both: 430 human genes carried a significant mutually exclusive isoform pair, 21 candidates surviving the bidirectional contrast filter, and 46 mouse candidates passed at the lower threshold. The strongest calls—TPM₁ in the cell line mixture, *Meg3* and *Malat1* in the mouse cortex—are consistent with documented isoform biology [41, 40], and most survived the swap of clustering strategy, so the signal does not depend on a particular partitioning scheme.

Broader Implications

This work has implications across three domains:

For computational genomics: The success of contingency-table statistics at transcript resolution demonstrates that sparsity can be modelled directly rather than overcome (Section 2.2). This philosophy may extend to other single-cell challenges—allele-specific expression, chromatin accessibility peak co-occurrence, and multi-modal integration.

For isoform biology: Detecting DTU in 3'-biased data opens new analysis opportunities for existing large-scale datasets (e.g., Human Cell Atlas, Tabula Sapiens¹) that were previously closed to isoform-level questions. While sensitivity is limited compared to full-length protocols, the trade-off between cell numbers and coverage may favor droplet methods when statistical power from thousands of cells outweighs per-cell depth.

For COTAN framework development: This thesis validates that COTAN's core design principles, set out in Section 2.2, generalize beyond gene-level analysis. Future extensions to other single-cell modalities (ATAC-seq² peak co-occurrence, protein abundance correlations in CITE-seq³) appear feasible.

What This Work Leaves Behind

Beyond the findings, this thesis leaves a set of reusable artefacts, published together in the public repository that accompanies it.

An analysis library. The `cotanisoform` R package collects the whole analysis path into documented functions: reading a count matrix, adapting it to COTAN, running the COTAN steps, extracting DTU candidates, and comparing runs. It can be installed and used on any single-cell dataset, independently of the two datasets studied here, and its logging layer lets a long run report its own progress and failures in one consistent trace.

A quantification pipeline. The Nextflow pipeline (Section 3.2) automates the path from raw reads to a count matrix, offering both the gene-level and the transcript-level index. Its steps run separately, so a re-run from an existing index does not repeat the download.

A reproducible analysis layer. The driver scripts reproduce the complete analysis of a dataset from a single configuration file, one per dataset and feature level, so each reported table can be regenerated without interactive input.

Documentation and reference outputs. The repository documents the method, the parameters behind each published table, and a parity script that re-derives the reported DTU tables from the stored objects, so a reader can check the results rather than take them on trust.

¹ **Human Cell Atlas / Tabula Sapiens:** Large international reference atlases of single-cell transcriptomes across human tissues, relied on 3'-biased droplet data and were therefore closed to isoform-level questions.

² **ATAC-seq:** Assay for Transposase-Accessible Chromatin: maps genome regions where chromatin is open, indicating regulatory activity.

³ **CITE-seq:** A protocol that simultaneously measures RNA and surface proteins in the same single cells, enabling multi-modal analysis.

Limitations and Future Work

The constraints on these results—3'-bias invisibility, the absence of orthogonal validation, the heuristic threshold τ , and EM rounding—are set out together with the directions they open in Section 4.4. The same underlying trade-off governs all four: sensitivity is bought with cells rather than coverage, so the value of the method grows as datasets do.

On the computational side, the pipeline and the analysis remain two halves joined at the count matrix: the Nextflow workflow ends at the filtered matrix, and the R analysis is started separately. Folding the analysis into the workflow as one further step would let a single run cover the path from raw accessions to DTU candidates, which is the main engineering direction this work leaves open.

A newer quantifier points the same way. SCALPEL[42] builds the transcript matrix from standard 3'-based data by a different route from the EM step used here, and its authors report higher sensitivity and closer recovery of true isoform proportions than existing tools, with experimental validation. Substituting it for `parsimony-em` would test whether the downstream DTU calls sharpen, making it a natural extension of this framework.

Final Remarks

This thesis demonstrates that transcript-level COTAN can reveal regulatory programs hidden to gene-centric analyses, and that a contingency-table approach to sparse data respects the measurement process rather than imposing transformations designed for dense, normally distributed data. The framework established here—methodological rigour, computational reproducibility and alignment with known isoform switches—provides a foundation for that future work, which still requires orthogonal validation and the complementary technologies set out in Section 4.4. Whether applied to neurogenesis, cancer evolution, or developmental biology, the core question remains: *which transcripts are cells expressing, and how does that choice change across conditions?*

Bibliography

- [1] Silvia Giulia Galfrè, Francesco Morandin, Matteo Pietrosanto, Federico Cremisi, and Manuela Helmer-Citterich. COTAN: scrna-seq data analysis based on gene co-expression. *NAR Genomics and Bioinformatics*, 3(3):lqab072, 2021. iii, 4, 7, 13, 16, 17, 18, 19, 27, 29
- [2] Matilde I. Conte, Azahara Fuentes-Trillo, and Cecilia Domínguez Conde. Opportunities and tradeoffs in single-cell transcriptomic technologies. *Trends in Genetics*, 40(1):83–93, 2024. iii, 3, 9
- [3] Jeroen Gilis, Kristoffer Vitting-Seerup, Koen Van den Berge, and Lieven Clement. satuRn: scalable analysis of differential transcript usage for bulk and single-cell RNA-sequencing applications. *F1000Research*, 10:374, 2021. iii, 4, 20
- [4] Ralph Patrick, David T. Humphreys, Vaibhao Janbandhu, Alicia Oshlack, Joshua W. K. Ho, Richard P. Harvey, and Kitty K. Lo. Sierra: discovery of differential transcript usage from poly-a-captured single-cell RNA-seq data. *Genome Biology*, 21:167, 2020. iii, 4
- [5] Dongze He, Mohsen Zakeri, Hirak Sarkar, Charlotte Soneson, Avi Srivastava, and Rob Patro. Alevin-fry unlocks rapid, accurate and memory-frugal quantification of single-cell RNA-seq data. *Nature Methods*, 19(3):316–322, 2022. iii, 11, 12, 14, 23, 25
- [6] A. Joglekar, C. Foord, J. Jarroux, S. Pollard, and H. U. Tilgner. From words to complete phrases: insight into single-cell isoforms using short and long reads. *Transcription*, 14(3–5):92–104, 2023. iii, 3
- [7] Claire Regan and Jonathan Preall. Practical considerations for single-cell genomics. *Current Protocols*, 2(8):e498, 2022. vii
- [8] Amber M. Young, Scott Van Buren, and Naim U. Rashid. Differential transcript usage analysis incorporating quantification uncertainty via compositional measurement error regression modeling. *Biostatistics*, 25(2):559–576, 2024. viii
- [9] Kristoffer Vitting-Seerup and Albin Sandelin. IsoformSwitchAnalyzeR: analysis of changes in genome-wide patterns of alternative splicing and its functional consequences. *Bioinformatics*, 33(17):2839–2841, 2017. viii
- [10] Randall K. Saiki, Stephen Scharf, Fred Faloona, Kary B. Mullis, Glenn T. Horn, Henry A. Erlich, and Norman Arnheim. Enzymatic amplification of beta-globin

genomic sequences and restriction site analysis for diagnosis of sickle cell anemia. *Science*, 230(4732):1350–1354, 1985. ix

- [11] Marina Cerqua, Valentina D’Addazio, Federica De Paolis, Maria Teresa Rossi, Ambra Turchetti, Emily Capone, Martina Fascione, Dalia Fuja, Rebecca Bellis, et al. MET δ 14 promotes a ligand-dependent, AKT-driven invasive growth. *Life Science Alliance*, 5(10):e202101325, 2022. 1
- [12] Fuchou Tang, Catalin Barbacioru, Yangzhou Wang, Ellen Nordman, Clarence Lee, Nanlan Xu, Xiaohui Wang, John Bode, Timothy Buck, Stuart Rosen, et al. mrna-seq whole-transcriptome analysis of a single cell. *Nature Methods*, 6(5):377–382, 2009. 2
- [13] Tamar Hashimshony, Florian Wagner, Noa Sher, and Itai Yanai. CEL-Seq: single-cell RNA-Seq by multiplexed linear amplification. *Cell Reports*, 2(3):666–673, 2012. 2
- [14] Daniel Ramsköld, Shujun Luo, Yu-Chieh Wang, Robin Li, Qiaolin Deng, Omid R Faridani, Gregory A Daniels, Irina Khrebtukova, Jeanne F Loring, Louise C Laurent, et al. Full-length mRNA-Seq from single-cell levels of RNA and individual circulating tumor cells. *Nature Biotechnology*, 30(8):777–782, 2012. 2
- [15] Simone Picelli, Åsa K Björklund, Omid R Faridani, Sven Sagasser, Gösta Winberg, and Rickard Sandberg. Smart-seq2 for sensitive full-length transcriptome profiling in single cells. *Nature Methods*, 10(11):1096–1098, 2013. 2
- [16] Grace X Y Zheng, Jessica M Terry, Phillip Belgrader, Paul Ryvkin, Zachary W Bent, Ryan Wilson, Solongo B Ziraldo, Tobias D Wheeler, Geoff P McDermott, Junjie Zhu, et al. Massively parallel digital transcriptional profiling of single cells. *Nature Communications*, 8:14049, 2017. 3, 23, 24
- [17] Shihao Shen, Juw Won Park, Zhi-xiang Lu, Lan Lin, Michael D Henry, Ying Nian Wu, Qing Zhou, and Yi Xing. rMATS: robust and flexible detection of differential alternative splicing from replicate RNA-Seq data. *Proceedings of the National Academy of Sciences*, 111(51):E5593–E5601, 2014. 4
- [18] Juan L Trincado, Juan C Entizne, Gerald Hysenaj, Babita Singh, Miha Skalic, David J Elliott, and Eduardo Eyras. SUPPA2: fast, accurate, and uncertainty-aware differential splicing analysis across multiple conditions. *Genome Biology*, 19(1):40, 2018. 4
- [19] Malgorzata Nowicka and Mark D Robinson. DRIMSeq: a dirichlet-multinomial framework for multivariate count outcomes in genomics. *F1000Research*, 5:1356, 2016. 4
- [20] Simon Anders, Alejandro Reyes, and Wolfgang Huber. Detecting differential usage of exons from RNA-seq data. *Genome Research*, 22(10):2008–2017, 2012. 4
- [21] Tobias Tekath and Martin Dugas. DTUrtle: a versatile pipeline for differential transcript usage analysis of bulk and single-cell RNA-seq data. *Bioinformatics*, 38(4):1133–1135, 2021. 4, 20

- [22] Jiarui Ding, Xian Adiconis, Sean K. Simmons, Monika S. Kowalczyk, Cynthia C. Hession, Nemanja D. Marjanovic, Travis K. Hughes, Marc H. Wadsworth, Tyler Burks, Lan T. Nguyen, et al. Systematic comparison of single-cell and single-nucleus RNA-sequencing methods. *Nature Biotechnology*, 38(6):737–746, 2020. 9, 24, 27
- [23] Ruochen Jiang, Tianyi Sun, Dongyuan Song, and Jingyi Jessica Li. Statistics or biology: the zero-inflation controversy about scrna-seq data. *Genome Biology*, 23(1):31, 2022. 9
- [24] Heng Li and Richard Durbin. Fast and accurate short read alignment with Burrows–Wheeler transform. *Bioinformatics*, 25(14):1754–1760, 2009. 11
- [25] Ben Langmead and Steven L. Salzberg. Fast gapped-read alignment with Bowtie 2. *Nature Methods*, 9(4):357–359, 2012. 11
- [26] Rob Patro, Geet Duggal, Michael I. Love, Rafael A. Irizarry, and Carl Kingsford. Salmon provides fast and bias-aware quantification of transcript expression. *Nature Methods*, 14(4):417–419, 2017. 11, 13, 23
- [27] Nicolas L. Bray, Harold Pimentel, Páll Melsted, and Lior Pachter. Near-optimal probabilistic RNA-seq quantification. *Nature Biotechnology*, 34(5):525–527, 2016. 11, 13
- [28] Arthur P. Dempster, Nan M. Laird, and Donald B. Rubin. Maximum likelihood from incomplete data via the EM algorithm. *Journal of the Royal Statistical Society: Series B*, 39(1):1–38, 1977. 14, 24
- [29] Michael I. Love, Charlotte Soneson, Rob Patro, and Mark D. Robinson. Swimming downstream: statistical analysis of differential transcript usage following Salmon quantification. *F1000Research*, 7:952, 2018. 14, 39
- [30] S. G. Galfrè, M. Fantozzi, A. Sîrbu, I. Testa, M. Tollosso, A. Alberti, C. Priami, and F. Morandin. From gene correlations to cell clusters: COTAN improved scrna-seq analysis. *NAR Genomics and Bioinformatics*, 2026. 16
- [31] Abhishek Sarkar and Matthew Stephens. Separating measurement and expression models clarifies confusion in single-cell RNA sequencing analysis. *Nature Genetics*, 53(6):770–777, 2021. 17
- [32] B.J. Draper, M.J. Dunning, and D.C. James. Selecting differential splicing methods: Practical considerations for short-read rna sequencing. *F1000Research*, 14:47, 2025. 20, 39
- [33] Tanya Barrett, Stephen E. Wilhite, Pierre Ledoux, Carlos Evangelista, Irene F. Kim, Maxim Tomashevsky, Kimberly A. Marshall, Katherine H. Phillippy, Patti M. Sherman, Michelle Holko, et al. NCBI GEO: archive for functional genomics data sets—update. *Nucleic Acids Research*, 41(D1):D991–D995, 2013. 23

- [34] Yuhao Hao, Tim Stuart, Madeline H. Kowalski, Saket Choudhary, Paul Hoffman, Austin Hartman, Avi Srivastava, Gesmira Molla, Shaista Madad, Carlos Fernandez-Granda, and Rahul Satija. Dictionary learning for integrative, multimodal and scalable single-cell analysis. *Nature Biotechnology*, 42(2):293–304, 2024. 23
- [35] Philip A. Ewels, Alexander Peltzer, Sven Fillinger, Harshil Patel, Johannes Alneberg, Andreas Wilm, Maxime Ulysse Garcia, Paolo Di Tommaso, and Sven Nahnsen. The nf-core framework for community-curated bioinformatics pipelines. *Nature Biotechnology*, 38(3):276–278, 2020. 23
- [36] Maddalena Arigoni, Maria Luisa Ratto, Federica Riccardo, Elisa Balmas, Lorenzo Calogero, Francesca Cordero, Marco Beccuti, Raffaele A. Calogero, and Luca Alessandri. A single cell RNAseq benchmark experiment embedding “controlled” cancer heterogeneity. *Scientific Data*, 11(1):159, 2024. 24, 27
- [37] Dongze He and Rob Patro. simpleaf: a simple, flexible, and scalable framework for single-cell data processing using alevin-fry. *Bioinformatics*, 39(10):btad614, 2023. 24
- [38] Paolo Di Tommaso, Maria Chatzou, Evan W. Floden, Pablo Prieto Barja, Emilio Palumbo, and Cedric Notredame. Nextflow enables reproducible computational workflows. *Nature Biotechnology*, 35(4):316–319, 2017. 26
- [39] Rasko Leinonen, Hideaki Sugawara, and Martin Shumway. The Sequence Read Archive. *Nucleic Acids Research*, 39(D1):D19–D21, 2011. 26
- [40] Shifeng Xue and Maria Barna. Specialized ribosomes: a new frontier in gene regulation and organismal biology. *Nature Reviews Molecular Cell Biology*, 13(6):355–369, 2012. 35, 41
- [41] Amr R. Ghanam, Shengwei Ke, Shujuan Wang, Ramy Elgendy, Chenyao Xie, Siqi Wang, Ran Zhang, Min Wei, Weiguang Liu, Jun Cao, Yan Zhang, Zhi Zhang, Tian Xue, Yong Zheng, and Xiaoyuan Song. Alternative transcribed 3’ isoform of long non-coding rna Malat1 inhibits mouse retinal oxidative stress. *iScience*, 26(1):105740, 2023. 35, 41
- [42] Franz Ake, Marcel Schilling, Sandra M. Fernández-Moya, Akshay Jaya Ganesh, Ana Gutiérrez-Franco, Lei Li, and Mireya Plass. Quantification of transcript isoforms at the single-cell level using SCALPEL. *Nature Communications*, 16(1):6402, 2025. 43

Appendix A

Supplementary Data

Complete DTU Results (Arigoni and Ding Datasets)

The tables below list every reported event. The same data are available as CSV in the public repository, so a reader can re-derive or re-filter them.

Table A.1: Significant DTU events in the human dataset (Arigoni 2024), including cluster enrichment and maximum contrasts.

Gene	Tr. A	Tr. B	Clusters A	Clusters B	ΔA	ΔB	COEX	p -value
TPM1	ENST00000334895	ENST00000358278	DV90, HCC78, HTB178	A549, CCL-185-IG	0.33	0.61	-0.101	7.72×10^{-67}
RPL10	ENST00000406022	ENST00000458500	A549	HCC78, PC9	0.22	0.46	-0.086	2.78×10^{-49}
RPL10	ENST00000458500	ENST00000491035	HCC78, PC9	A549, PBMCs	0.46	0.25	-0.084	6.41×10^{-47}
TPM1	ENST00000404484	ENST00000651344	HCC78, HTB178	A549, CCL-185-IG	0.27	0.37	-0.074	1.15×10^{-36}
TPM1	ENST00000317516	ENST00000651344	HCC78, HTB178	A549, CCL-185-IG	0.28	0.36	-0.068	4.41×10^{-31}
RPS24	ENST00000360830	ENST00000465692	DV90, HCC78	CCL-185-IG, HTB178, PBMCs	0.28	0.30	-0.051	1.76×10^{-18}
TPM1	ENST00000288398	ENST00000334895	A549, CCL-185-IG	HTB178	0.30	0.21	-0.047	1.13×10^{-15}
TPM1	ENST00000560959	ENST00000651344	HCC78, HTB178	A549, CCL-185-IG	0.22	0.32	-0.044	8.17×10^{-14}
RPS24	ENST00000360830	ENST00000482069	DV90, HCC78	CCL-185-IG, HTB178, PBMCs	0.28	0.28	-0.040	8.15×10^{-12}
TPM1	ENST00000334895	ENST00000651344	DV90, HCC78, HTB178	A549, CCL-185-IG	0.28	0.44	-0.040	8.76×10^{-12}
DSTN	ENST00000449141	ENST00000474024	HTB178	CCL-185-IG	0.28	0.22	-0.039	3.70×10^{-11}
SAT1	ENST00000379253	ENST00000462639	A549	PC9	0.26	0.27	-0.038	6.62×10^{-11}
TPM4	ENST00000645471	ENST00000646974	CCL-185-IG	HCC78	0.32	0.39	-0.034	5.27×10^{-09}
SAT1	ENST00000379251	ENST00000462639	A549	PC9	0.22	0.21	-0.031	7.47×10^{-08}
SAT1	ENST00000379270	ENST00000487713	A549	PC9	0.25	0.25	-0.031	9.35×10^{-08}
RPL35A	ENST00000429437	ENST00000642387	PC9	PBMCs	0.21	0.22	-0.031	1.26×10^{-07}
TPM1	ENST00000558544	ENST00000651344	HCC78, HTB178	A549, CCL-185-IG	0.23	0.33	-0.019	1.04×10^{-03}
TPM1	ENST00000358278	ENST00000558868	A549, CCL-185-IG	CRL5868, DV90, HCC78, HTB178	0.52	0.27	-0.018	1.96×10^{-03}
SNHG5	ENST00000654795	ENST00000664396	CCL-185-IG	PC9	0.23	0.28	-0.018	2.37×10^{-03}
RPS8	ENST00000372209	ENST00000485390	PC9	PBMCs	0.21	0.29	-0.017	4.19×10^{-03}
TPM1	ENST00000317516	ENST00000358278	DV90, HCC78, HTB178	A549, CCL-185-IG	0.33	0.54	-0.012	4.03×10^{-02}

Table A.2: Shared DTU events detected by both clustering methods (Ding dataset), including cell types and contrasts.

Gene	Tr. A	Tr. B	Clusters A	Clusters B	Δ A	Δ B	COEX	p -value
Meg3	ENSMUST00000146701	ENSMUST00000246047	07, 08, 09, 10, 11, 12	−1, 02, 03, 04, 05, 06	0.28	0.66	−0.174	1.58×10^{-28}
Meg3	ENSMUST00000129245	ENSMUST00000146701	−1, 02, 03, 04, 05, 06	07, 08, 09, 10, 11, 12	0.63	0.28	−0.160	1.65×10^{-24}
Malat1	ENSMUST00000173499	ENSMUST00000245771	08, 10, 11, 12	02, 04, 05, 06	0.27	0.29	−0.104	2.91×10^{-11}
Malat1	ENSMUST00000173499	ENSMUST00000245150	08, 09, 10, 11, 12	−1, 02, 04, 05, 06	0.17	0.29	−0.073	3.30×10^{-06}
Meg3	ENSMUST00000129245	ENSMUST00000247387	02, 04, 05, 06	11, 12	0.34	0.29	−0.043	5.70×10^{-03}
Meg3	ENSMUST00000129245	ENSMUST00000246553	02, 04, 05, 06	11, 12	0.34	0.29	−0.042	6.88×10^{-03}
Meg3	ENSMUST00000129245	ENSMUST00000245180	02, 04, 05, 06	11, 12	0.33	0.30	−0.042	7.27×10^{-03}
Meg3	ENSMUST00000129245	ENSMUST00000246499	02, 04, 05, 06	11, 12	0.33	0.29	−0.042	7.40×10^{-03}
Meg3	ENSMUST00000246047	ENSMUST00000247387	02, 04, 05, 06	11, 12	0.37	0.29	−0.042	7.96×10^{-03}
Meg3	ENSMUST00000129245	ENSMUST00000246427	02, 04, 05, 06	11, 12	0.33	0.29	−0.041	8.84×10^{-03}
Meg3	ENSMUST00000246047	ENSMUST00000246553	02, 04, 05, 06	11, 12	0.37	0.29	−0.041	9.49×10^{-03}
Meg3	ENSMUST00000129245	ENSMUST00000246700	02, 04, 05, 06	11, 12	0.33	0.29	−0.041	9.51×10^{-03}
Meg3	ENSMUST00000245180	ENSMUST00000246047	11, 12	02, 04, 05, 06	0.29	0.37	−0.041	9.66×10^{-03}
Meg3	ENSMUST00000246047	ENSMUST00000246499	02, 04, 05, 06	11, 12	0.36	0.29	−0.040	9.79×10^{-03}
Meg3	ENSMUST00000129245	ENSMUST00000143836	02, 04, 05, 06	11, 12	0.33	0.29	−0.040	9.83×10^{-03}
Meg3	ENSMUST00000129245	ENSMUST00000245980	02, 04, 05, 06	11, 12	0.33	0.29	−0.040	1.04×10^{-02}
Meg3	ENSMUST00000246047	ENSMUST00000246427	02, 04, 05, 06	11, 12	0.36	0.29	−0.040	1.15×10^{-02}
Meg3	ENSMUST00000246047	ENSMUST00000246700	02, 04, 05, 06	11, 12	0.36	0.29	−0.039	1.24×10^{-02}
Meg3	ENSMUST00000129245	ENSMUST00000247194	02, 04, 05, 06	11, 12	0.32	0.29	−0.039	1.26×10^{-02}
Meg3	ENSMUST00000143836	ENSMUST00000246047	11, 12	02, 04, 05, 06	0.28	0.36	−0.039	1.27×10^{-02}
Meg3	ENSMUST00000245980	ENSMUST00000246047	11, 12	02, 04, 05, 06	0.28	0.36	−0.039	1.36×10^{-02}
Meg3	ENSMUST00000246047	ENSMUST00000247194	02, 04, 05, 06	11, 12	0.36	0.28	−0.038	1.60×10^{-02}
Pde10a	ENSMUST00000141877	ENSMUST00000233652	07	09, 10, 11	0.51	0.13	−0.037	1.67×10^{-02}
Meg3	ENSMUST00000129245	ENSMUST00000243087	02, 04, 05, 06	11, 12	0.32	0.24	−0.034	3.21×10^{-02}
Grin2b	ENSMUST00000053880	ENSMUST00000125905	11	02, 04	0.10	0.10	−0.033	3.39×10^{-02}
Meg3	ENSMUST00000129245	ENSMUST00000246154	02, 04, 05, 06	11, 12	0.32	0.25	−0.033	3.47×10^{-02}
Meg3	ENSMUST00000129245	ENSMUST00000243937	02, 04, 05, 06	11, 12	0.31	0.25	−0.033	3.79×10^{-02}
Meg3	ENSMUST00000243087	ENSMUST00000246047	11, 12	02, 04, 05, 06	0.24	0.35	−0.032	3.82×10^{-02}
Meg3	ENSMUST00000129245	ENSMUST00000143847	02, 04, 05, 06	11, 12	0.31	0.24	−0.032	3.98×10^{-02}

Continued on next page

Gene	Tr. A	Tr. B	Clusters A	Clusters B	Δ A	Δ B	COEX	<i>p</i> -value
Meg3	ENSMUST00000129245	ENSMUST00000247320	02, 04, 05, 06	11, 12	0.31	0.24	−0.032	3.98×10^{-02}
Meg3	ENSMUST00000126289	ENSMUST00000129245	11, 12	02, 04, 05, 06	0.25	0.31	−0.032	4.08×10^{-02}
Meg3	ENSMUST00000246047	ENSMUST00000246154	02, 04, 05, 06	11, 12	0.35	0.24	−0.032	4.12×10^{-02}
Meg3	ENSMUST00000129245	ENSMUST00000247335	02, 04, 05, 06	11, 12	0.31	0.25	−0.032	4.19×10^{-02}
Meg3	ENSMUST00000243937	ENSMUST00000246047	11, 12	02, 04, 05, 06	0.24	0.35	−0.031	4.45×10^{-02}
Meg3	ENSMUST00000143847	ENSMUST00000246047	11, 12	02, 04, 05, 06	0.24	0.35	−0.031	4.66×10^{-02}
Meg3	ENSMUST00000246047	ENSMUST00000247320	02, 04, 05, 06	11, 12	0.35	0.24	−0.031	4.66×10^{-02}
Meg3	ENSMUST00000129245	ENSMUST00000246462	02, 04, 05, 06	11, 12	0.31	0.24	−0.031	4.73×10^{-02}
Meg3	ENSMUST00000126289	ENSMUST00000246047	11, 12	02, 04, 05, 06	0.25	0.35	−0.031	4.81×10^{-02}
Meg3	ENSMUST00000246047	ENSMUST00000247335	02, 04, 05, 06	11, 12	0.35	0.24	−0.031	4.89×10^{-02}

Table A.3: DTU events exclusive to transcript-level clustering (Ding dataset).

Gene	Tr. A	Tr. B	Clusters A	Clusters B	Δ A	Δ B	COEX	p -value
Srrm2	ENSMUST00000088621	ENSMUST00000186961	10	09, 11	0.25	0.10	-0.059	1.53×10^{-04}
Srrm2	ENSMUST00000088621	ENSMUST00000233636	10	09, 11	0.26	0.12	-0.059	1.78×10^{-04}
Ttc14	ENSMUST00000117915	ENSMUST00000200271	-1	11	0.07	0.06	-0.036	2.19×10^{-02}

Table A.4: DTU events exclusive to gene-level clustering (Ding dataset).

Gene	Tr. A	Tr. B	Clusters A	Clusters B	Δ A	Δ B	COEX	p -value
Snrnp70	ENSMUST00000074575	ENSMUST00000211178	03	14	0.06	0.06	-0.042	6.81×10^{-03}
Lrrc7	ENSMUST00000197866	ENSMUST00000198284	05	08	0.06	0.06	-0.040	9.87×10^{-03}
Zfp941	ENSMUST00000080651	ENSMUST00000106052	12	04	0.08	0.07	-0.032	3.87×10^{-02}
Meg3	ENSMUST00000243087	ENSMUST00000247243	12	-1	0.08	0.06	-0.031	4.89×10^{-02}